



Onyx

Distributed Workflows for Dynamic Systems

Michael Drogalis

Me!

- Independent consultant
- Functional programming, Clojure, distributed systems
- @MichaelDrogalis



What is Onyx?

- A new kind of distributed computation system
- Provides an information model for the developer
- Inverts calling control of traditional frameworks
- Competes in stream, batch, and ingestion space
- Enabled by hardware advances in the last decade



Onyx Goals

- Take apart the monolith - data and fns
- Target Clojure & the JVM
- Batch and stream hybrid, transparent code reuse
- Transactional execution semantics
- Powerful extensibility API
- An API dedicated to side-effects & state

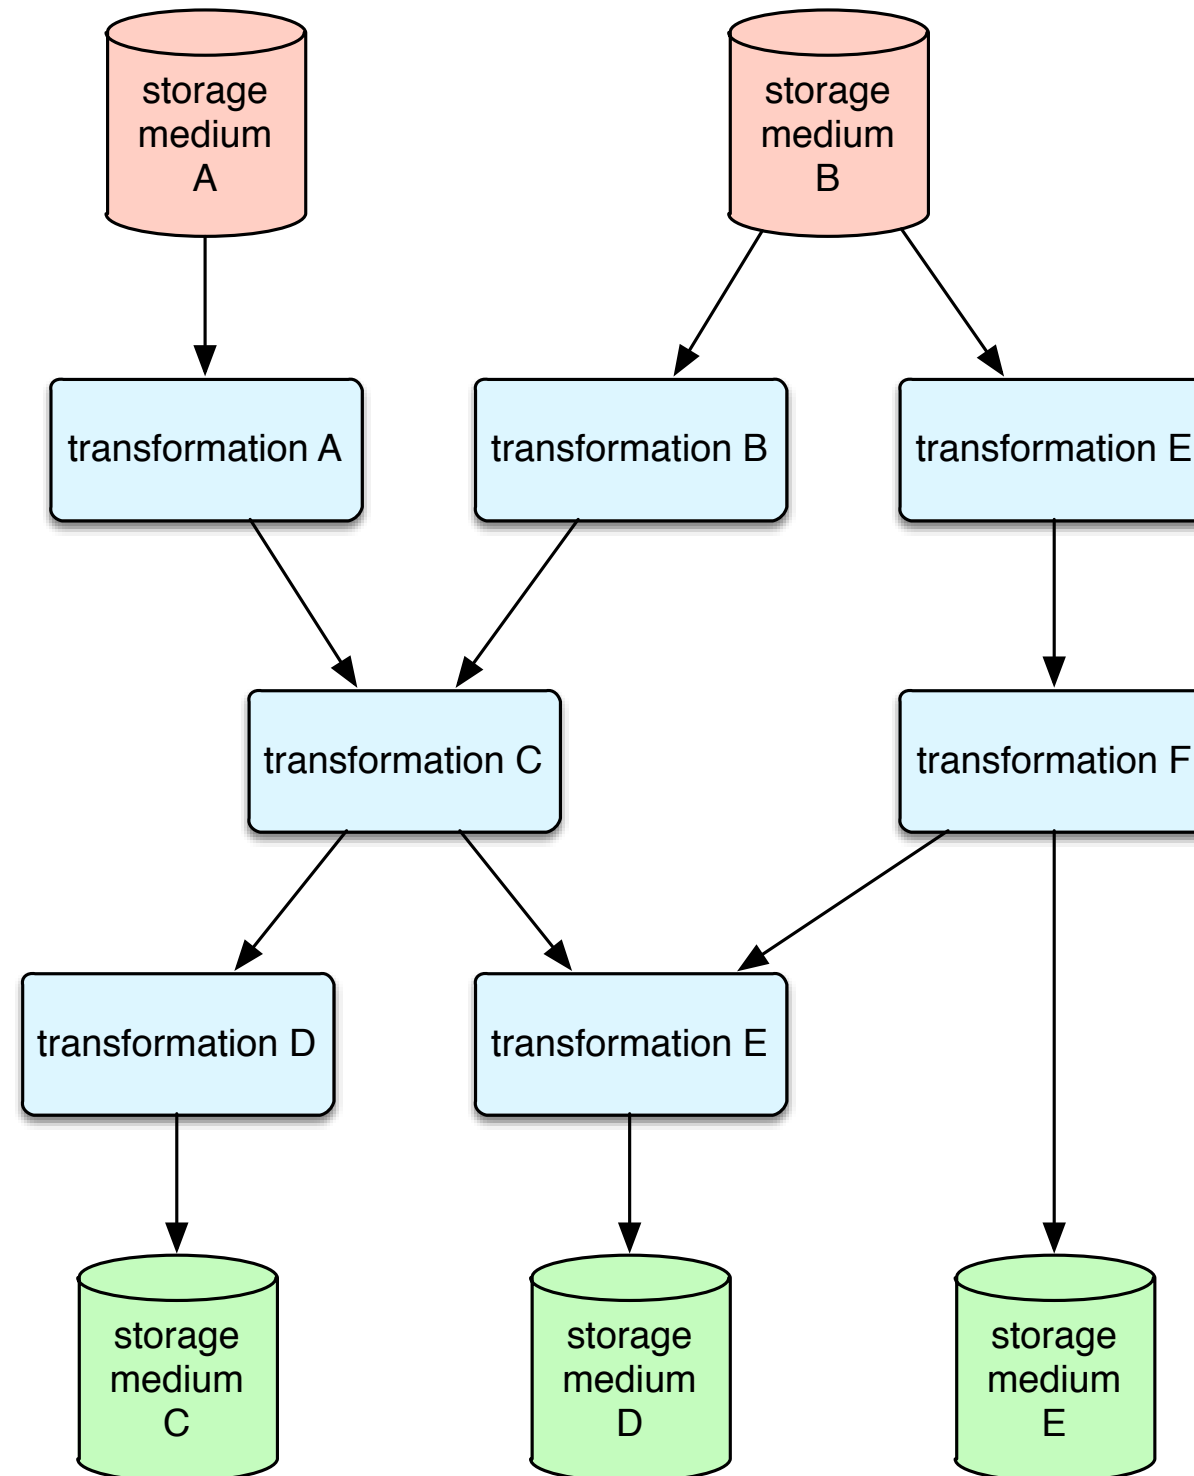


Overview

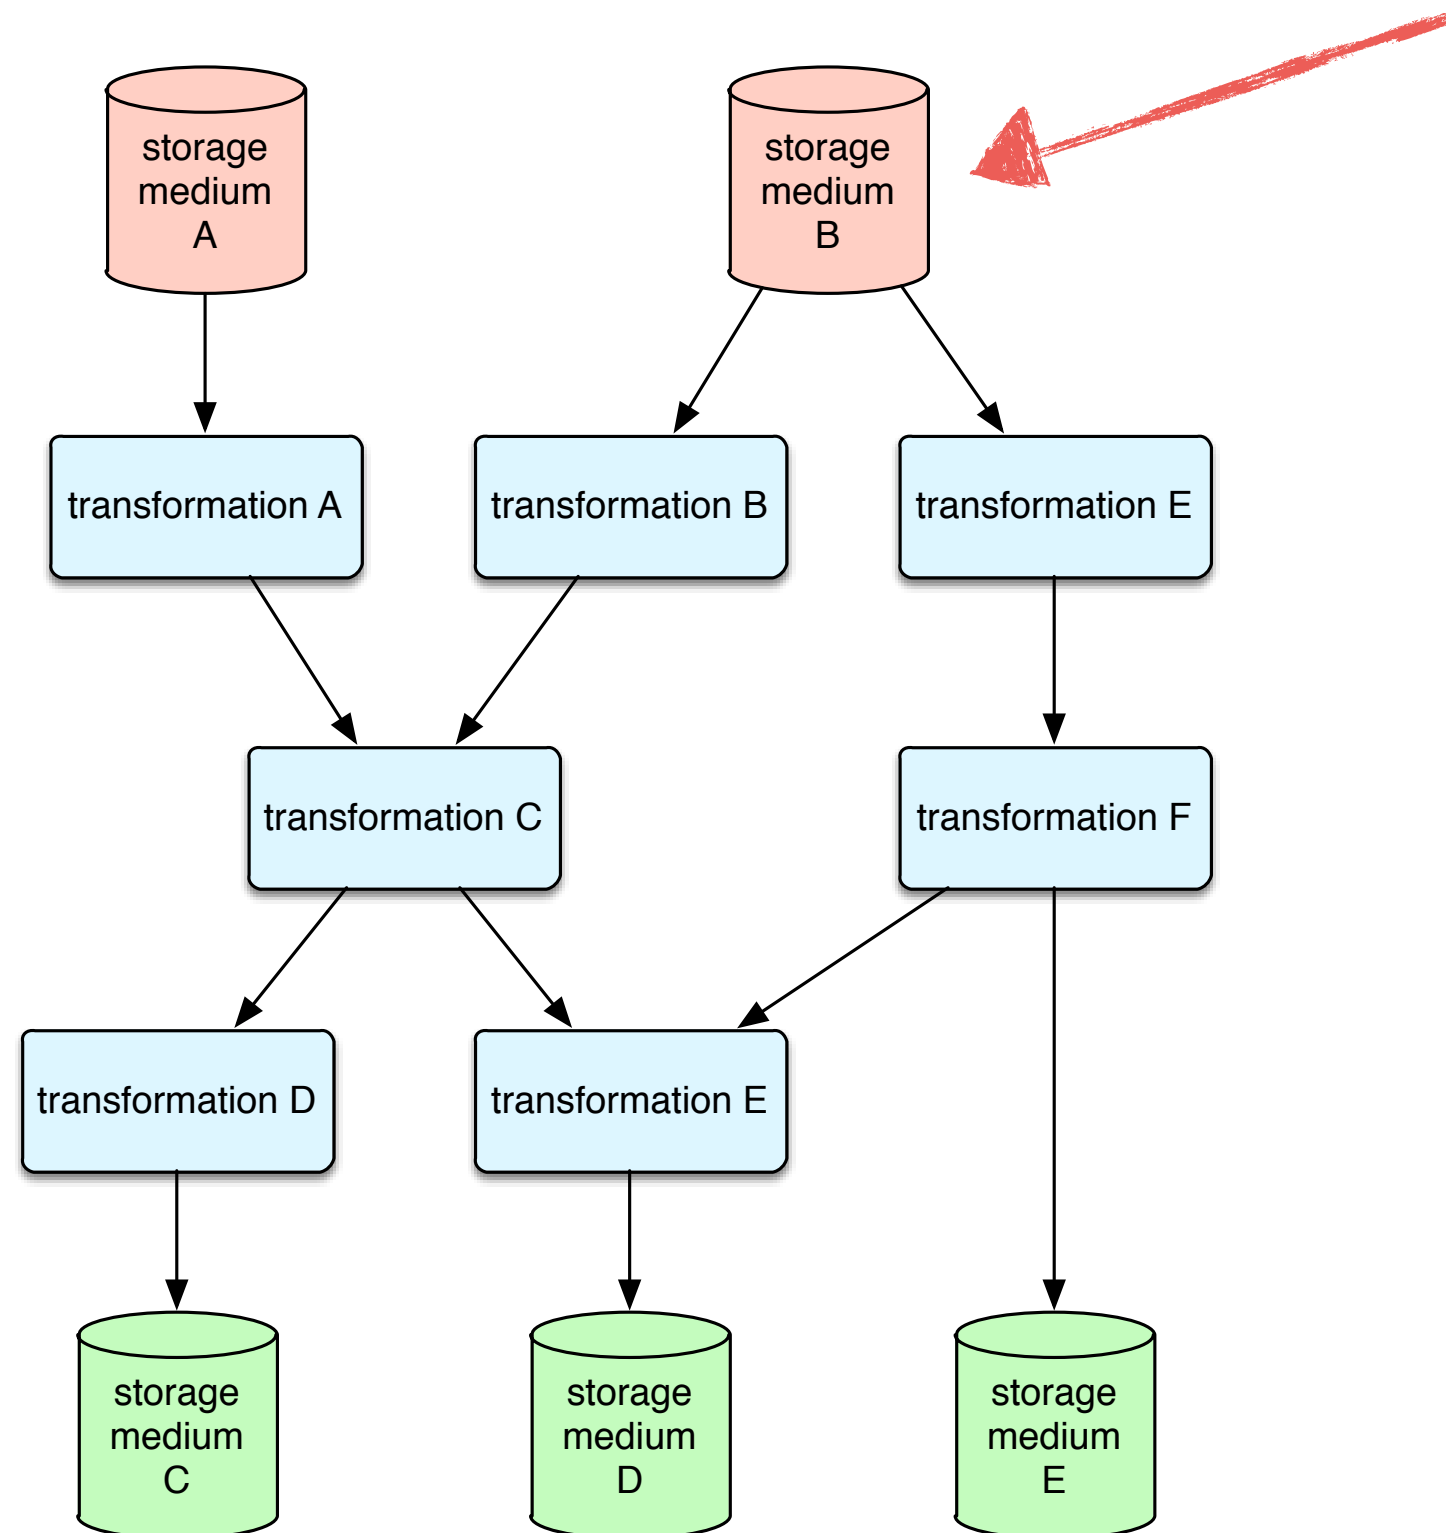
- The problem space
- Information model and APIs
- Architecture
- Development experience



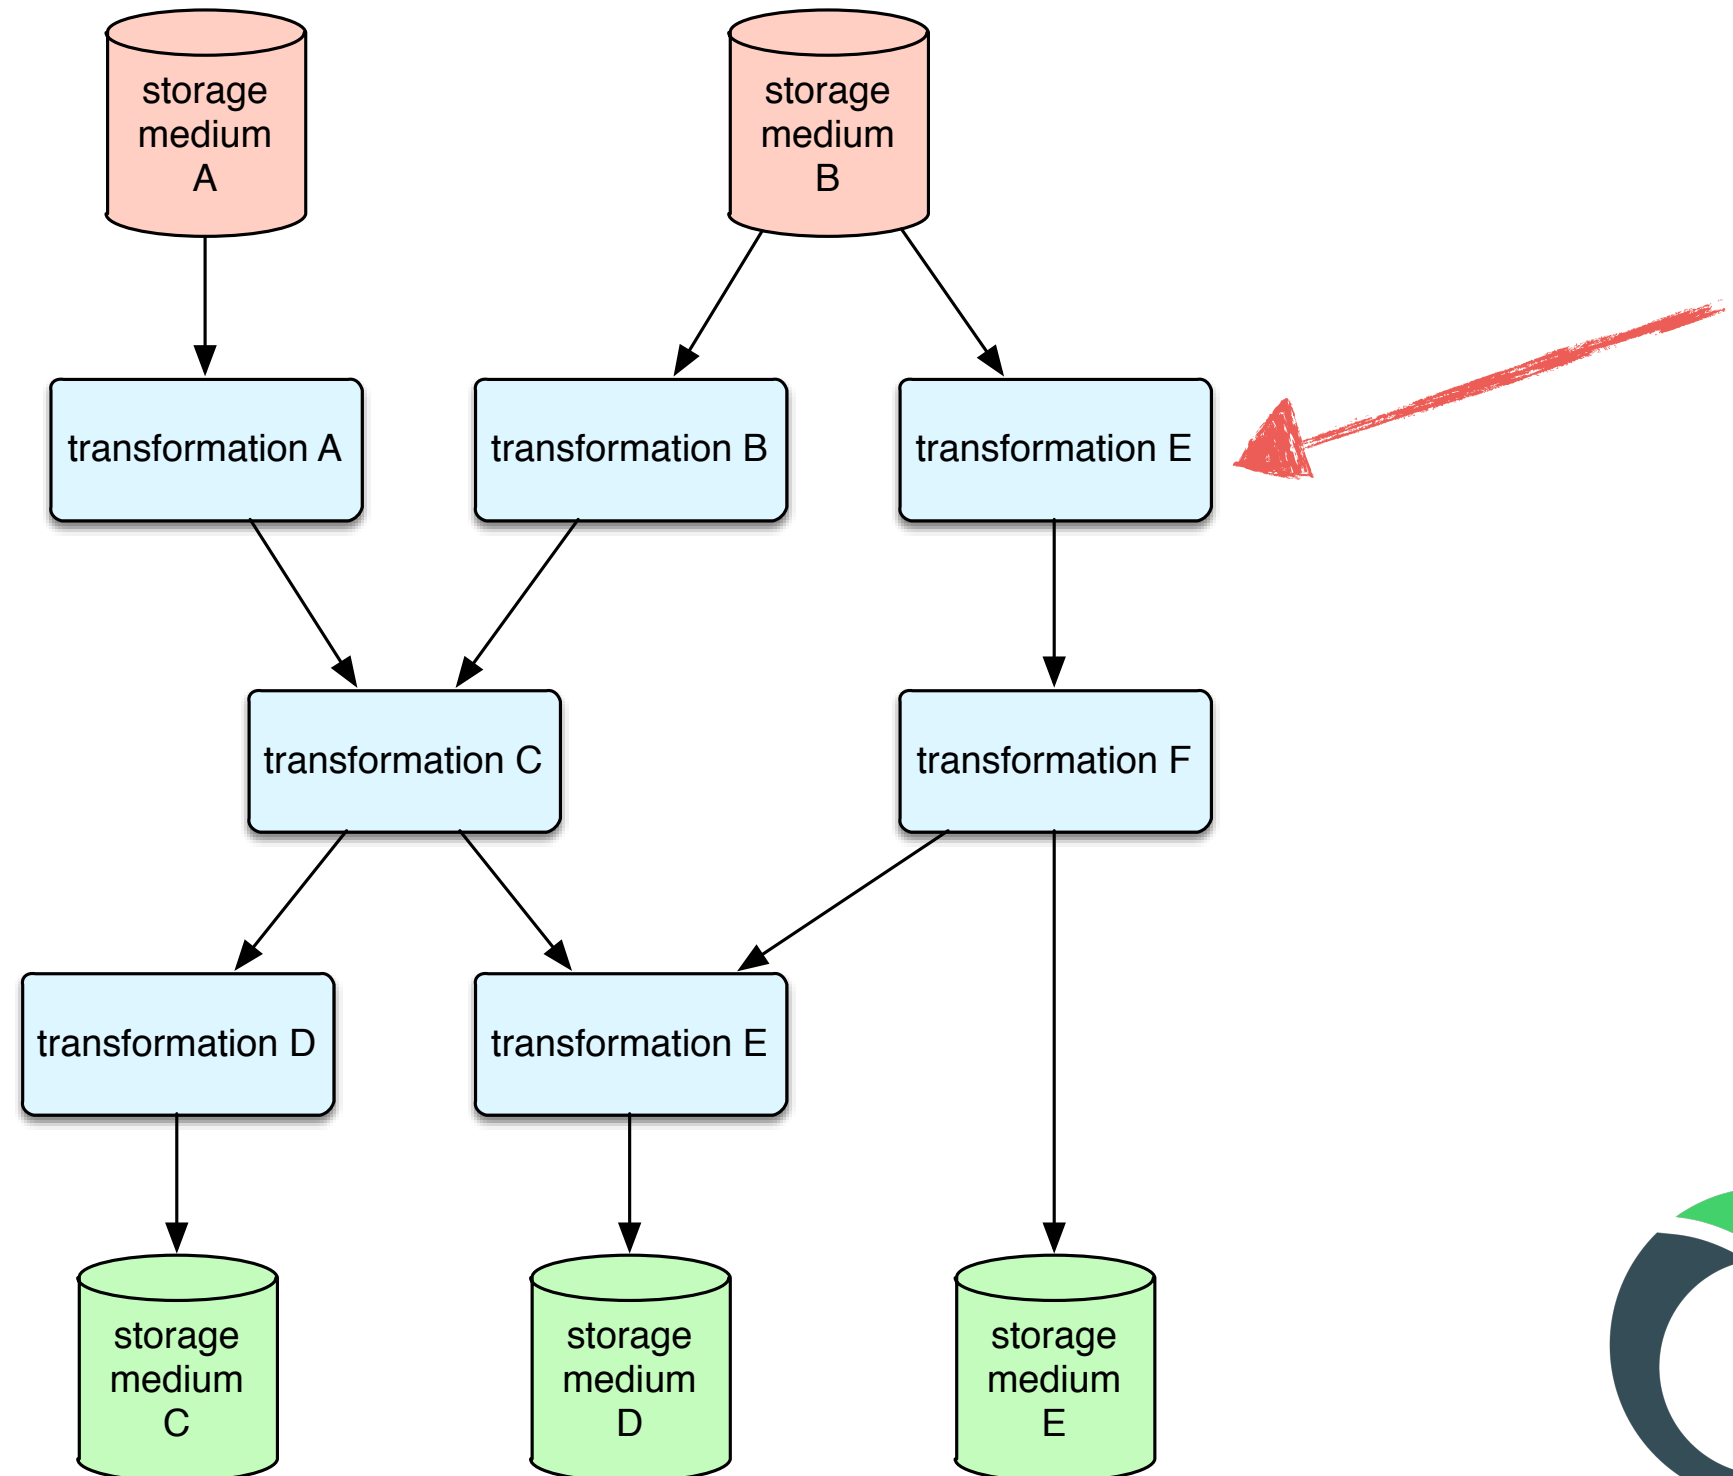
The Problem



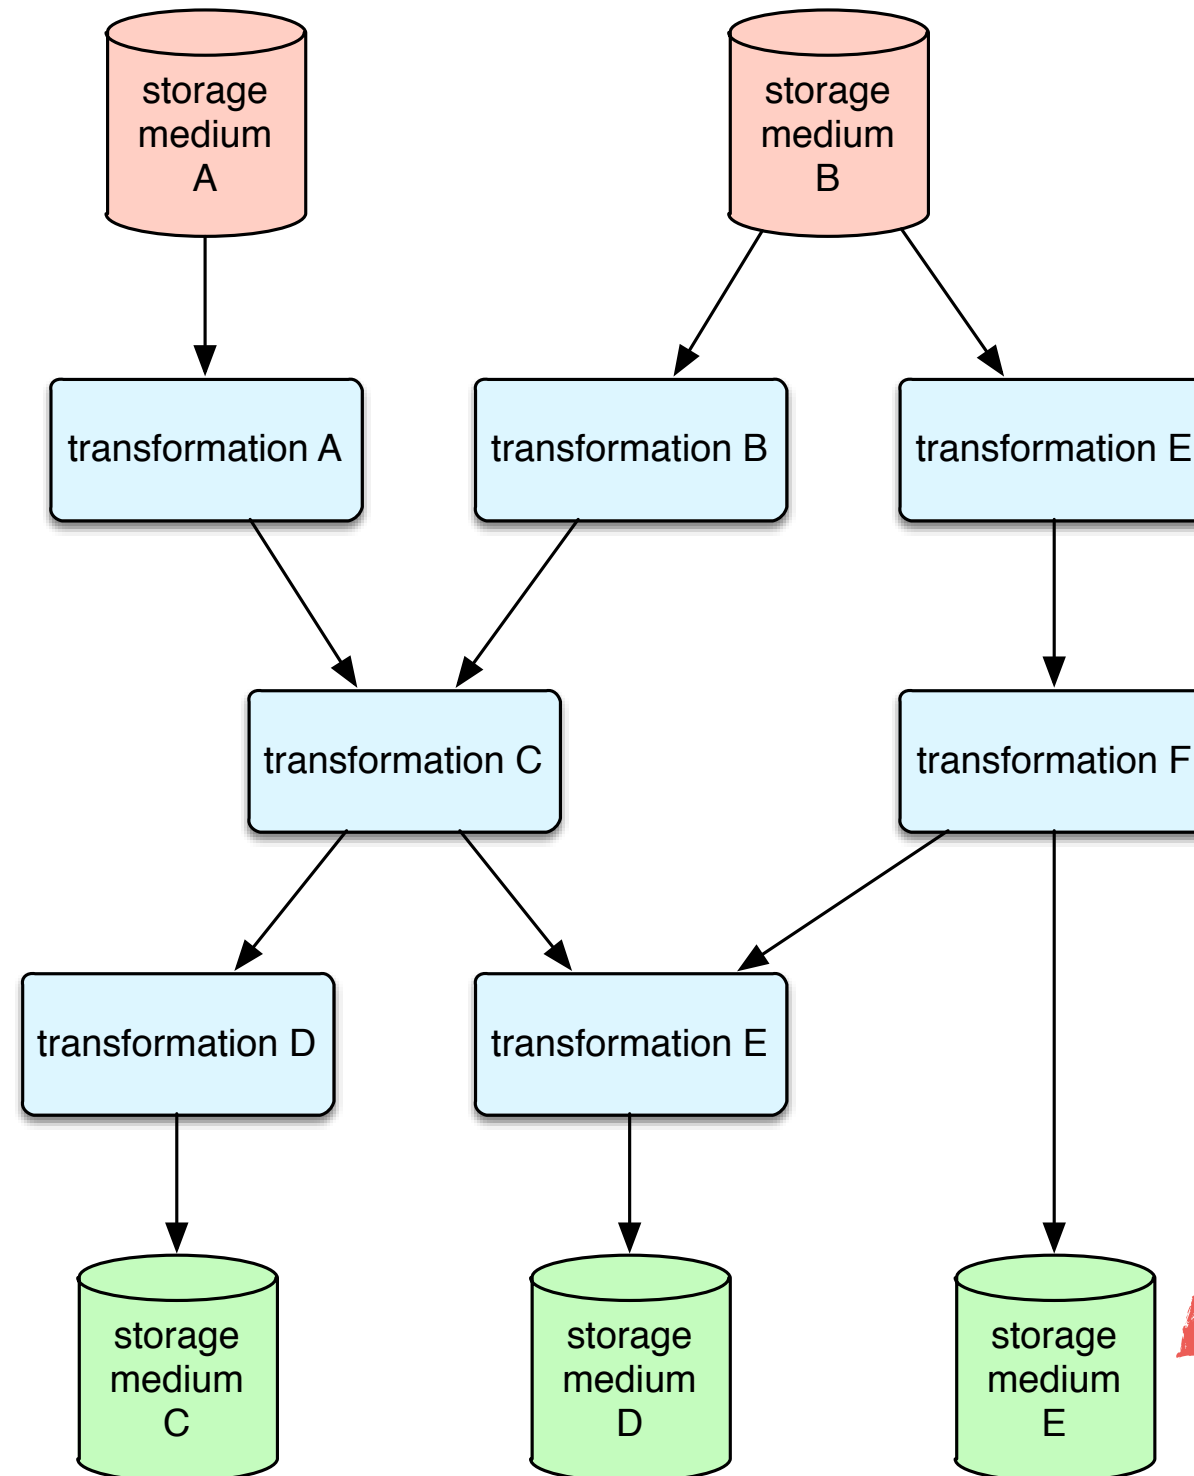
The Problem



The Problem



The Problem



The Problem

- Inputs, transformations, and outputs not static
- Determined at runtime - usually in another ecosystem!
- Specification of distributed execution needs to be...



The Problem

- Inputs, transformations, and outputs not static
- Determined at runtime - usually in another ecosystem!
- Specification of distributed execution needs to be...
 - language agnostic
 - location agnostic
 - temporally agnostic
 - tolerant to machine generation



Cascading: A Solution?

```
Pipe stopPipe = new Pipe("stop");  
Pipe tokenPipe = new HashJoin(docPipe, token, stopPipe, stop, new LeftJoin());  
  
tokenPipe = new Each(tokenPipe, stop, new RegexFilter("^$"));  
tokenPipe = new Retain(tokenPipe, fieldSelector);
```



Cascading: A Solution?

```
Pipe stopPipe = new Pipe("stop");  
Pipe tokenPipe = new HashJoin(docPipe, token, stopPipe, stop, new LeftJoin());  
  
tokenPipe = new Each(tokenPipe, stop, new RegexFilter("^$"));  
tokenPipe = new Retain(tokenPipe, fieldSelector);
```

~~language agnostic~~

~~location agnostic~~

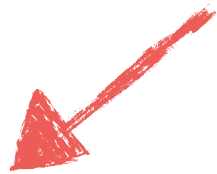
~~temporally agnostic~~

machine tolerant



Cascading: A Solution?

mechanism



```
Pipe stopPipe = new Pipe("stop");  
Pipe tokenPipe = new HashJoin(docPipe, token, stopPipe, stop, new LeftJoin());
```

```
tokenPipe = new Each(tokenPipe, stop, new RegexFilter("^$"));  
tokenPipe = new Retain(tokenPipe, fieldSelector);
```



~~language agnostic~~

~~location agnostic~~

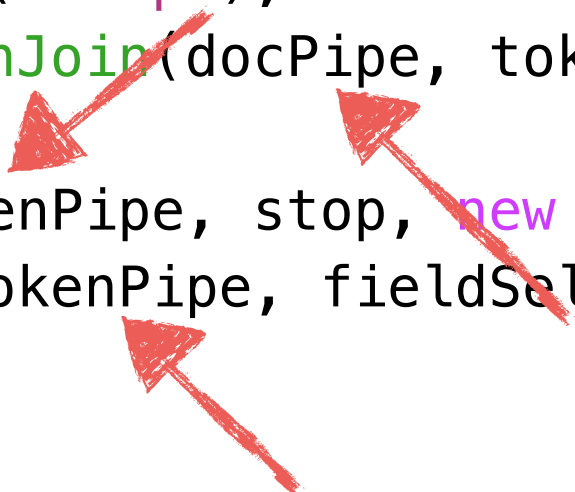
~~temporally agnostic~~

machine tolerant



Cascading: A Solution?

```
Pipe stopPipe = new Pipe("stop");  
Pipe tokenPipe = new HashJoin(docPipe, token, stopPipe, stop, new LeftJoin());  
  
tokenPipe = new Each(tokenPipe, stop, new RegexFilter("^$"));  
tokenPipe = new Retain(tokenPipe, fieldSelector);
```



structure

~~language agnostic~~

~~location agnostic~~



~~temporally agnostic~~

machine tolerant



Cascading: A Solution?

```
Pipe stopPipe = new Pipe("stop");  
Pipe tokenPipe = new HashJoin(docPipe, token, stopPipe, stop, new LeftJoin());  
  
tokenPipe = new Each(tokenPipe, stop, new RegexFilter("^$"));  
tokenPipe = new Retain(tokenPipe, fieldSelector);
```



configuration

~~language agnostic~~

~~location agnostic~~

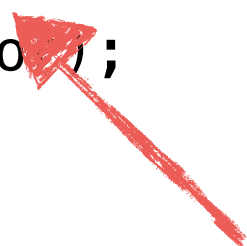

~~temporally agnostic~~

machine tolerant



Cascading: A Solution?

```
Pipe stopPipe = new Pipe("stop");  
Pipe tokenPipe = new HashJoin(docPipe, token, stopPipe, stop, new LeftJoin());  
  
tokenPipe = new Each(tokenPipe, stop, new RegexFilter("^$"));  
tokenPipe = new Retain(tokenPipe, fieldSelector);
```



forced concretion

~~language agnostic~~

~~location agnostic~~

~~temporally agnostic~~

machine tolerant



Cascading: A Solution?

```
Pipe stopPipe = new Pipe("stop");  
Pipe tokenPipe = new HashJoin(docPipe, token, stopPipe, stop, new LeftJoin());  
  
tokenPipe = new Each(tokenPipe, stop, new RegexFilter("^$"));  
tokenPipe = new Retain(tokenPipe, fieldSelector);
```



black box

~~language agnostic~~

~~location agnostic~~

~~temporally agnostic~~

machine tolerant



Storm: A Solution?

```
(topology
  {"1" (spout-spec sentence-spout)
    "2" (spout-spec other-spout :p 2)}
  {"3" (bolt-spec {"1" :shuffle "2" :shuffle}
    split-sentence :p 5)
    "4" (bolt-spec {"3" ["word"]}
    word-count :p 6))}
```



Storm: A Solution?

```
(topology
  {"1" (spout-spec sentence-spout)
    "2" (spout-spec other-spout :p 2)}
  {"3" (bolt-spec {"1" :shuffle "2" :shuffle}
    split-sentence :p 5)
    "4" (bolt-spec {"3" ["word"]}
    word-count :p 6))}
```

~~language agnostic~~

location agnostic

temporally agnostic

~~machine tolerant~~



Cascalog: A Solution?

```
(?- (stdout)
    (<- [?word ?count]
      (sentence :> ?line)
      (tokenize :< ?line :> ?word)
      (count :> ?count)))
```



Cascalog: A Solution?

```
(?- (stdout)
    (<- [?word ?count]
      (sentence :> ?line)
      (tokenize :< ?line :> ?word)
      (count :> ?count)))
```

~~language agnostic~~

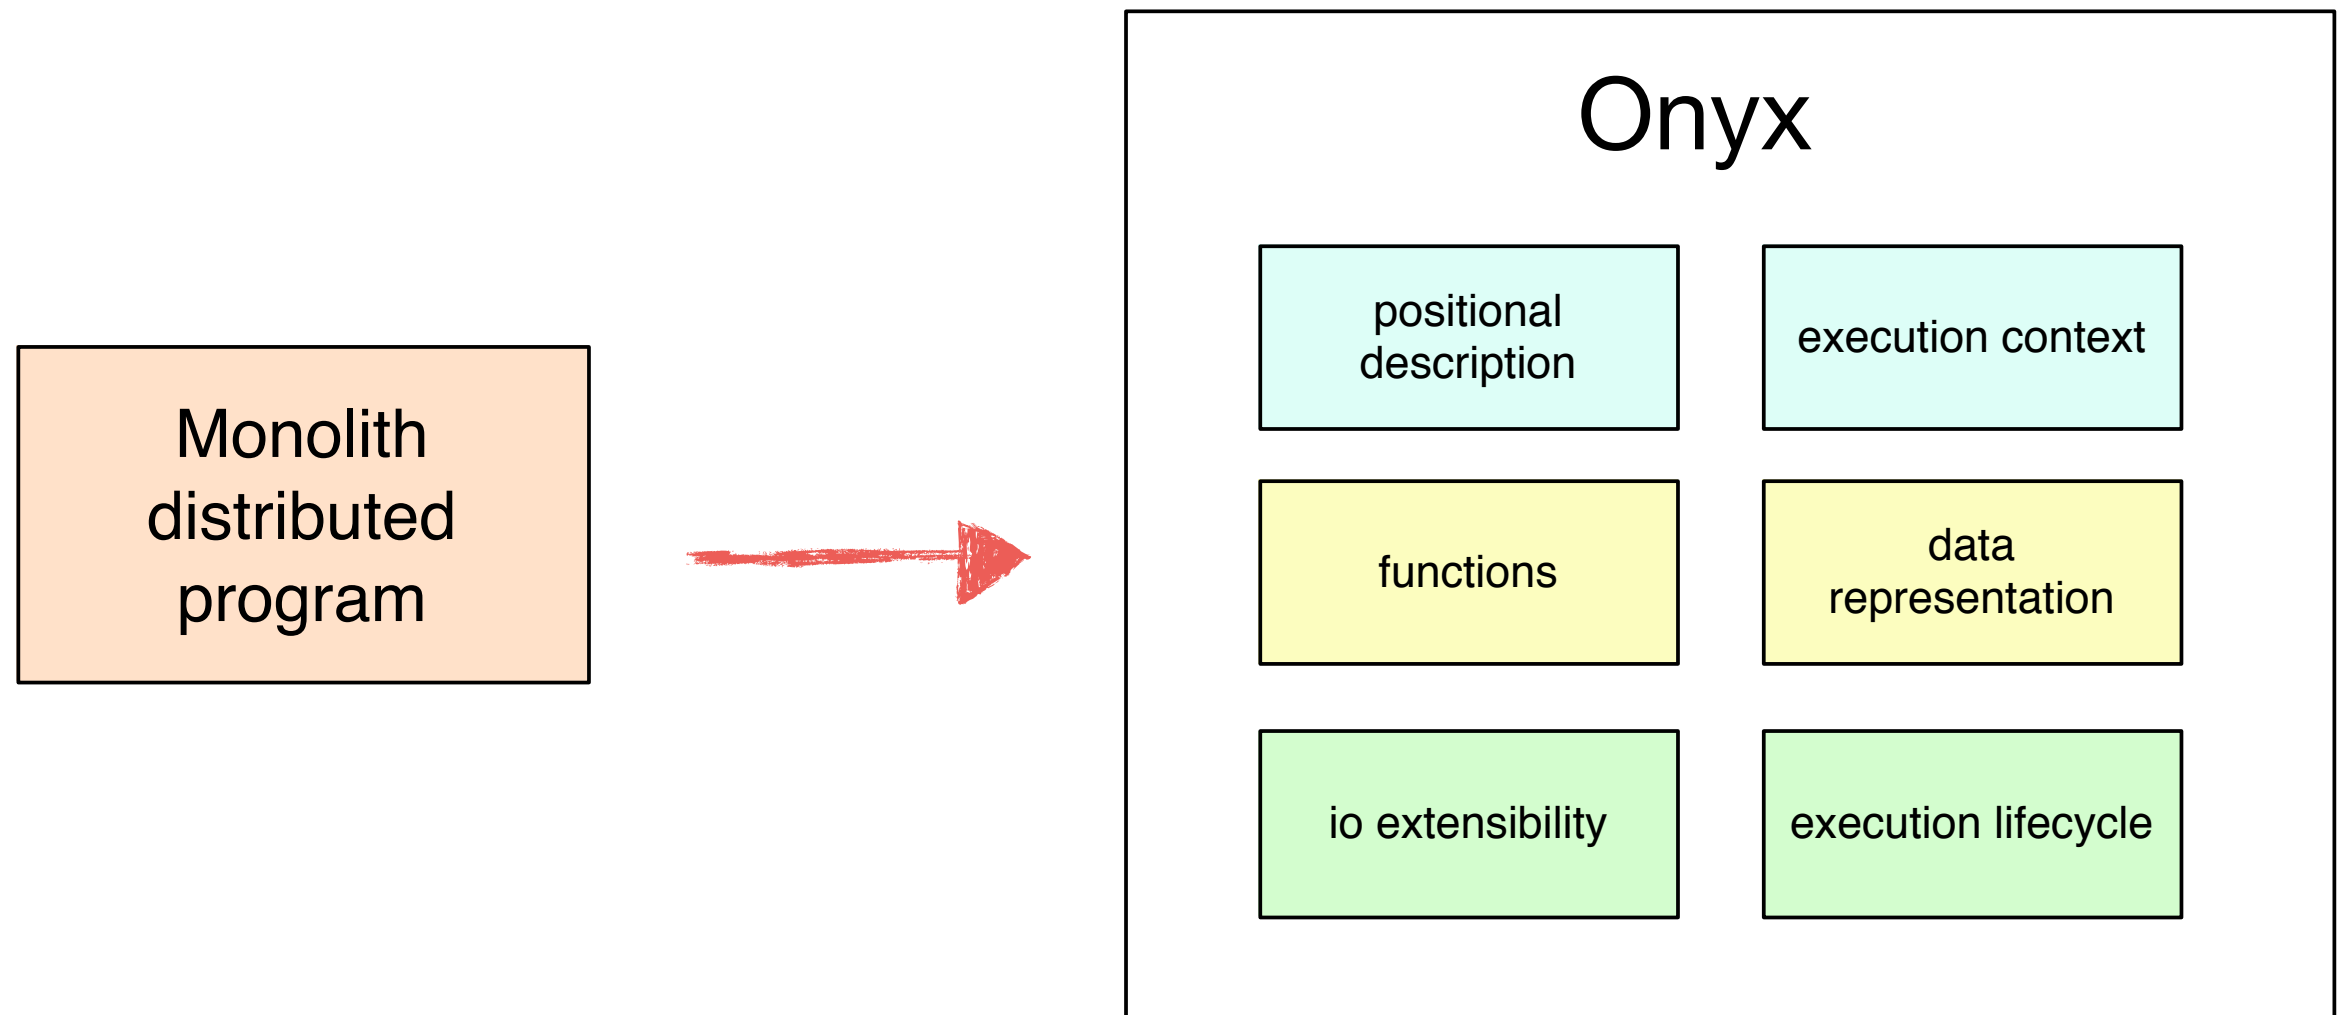
~~location agnostic~~

~~temporally agnostic~~

~~machine tolerant~~



Decomposing the Monolith



Guiding Example

name

age

Mike

23

John

19

Kristen

24

Guiding Example

name

age

Mike

23

John

19

Kristen

24

name

age

MIKE

24

JOHN

20

KRISTEN

25

Data Representation

- Declare names of values from fn -> fn
- Typical representation is tuple and fields
- Ordered, explicit sequence of values

```
(defbolt transform-person ["name" "age"] [tuple collector]  
  (emit-bolt! collector ["MIKE" 24])  
  (ack! collector tuple))
```

output fields

output tuple



Data Representation

- **Segment** - just a Clojure map
- No notion of ordering, no explicit declaration
- Intermediaries don't have to care
- Forces values to be named

```
{:name "Mike"  
  :age 24}
```



Functions

- Plain Clojure functions
- Input: a single segment (additional params allowed)
- Output: one segment, or a seq of segments
- No macros, calling context, stateful parameters

```
(defn transform-person [segment]
  (-> segment
    (update-in [:name] clojure.string/upper-case)
    (update-in [:age] inc)))
```



Did Somebody Say Transducers?!

```
(defn transform-person [segment]
  (-> segment
    (update-in [:name] clojure.string/upper-case)
    (update-in [:age] inc)))
```


```
(defn derek? [segment]
  (= (:name segment) "Derek"))
```

```
(defn x-form []
  (comp (remove derek?)
    (map transform-person)))
```

- Available in edge build (0.4.0-SNAPSHOT)

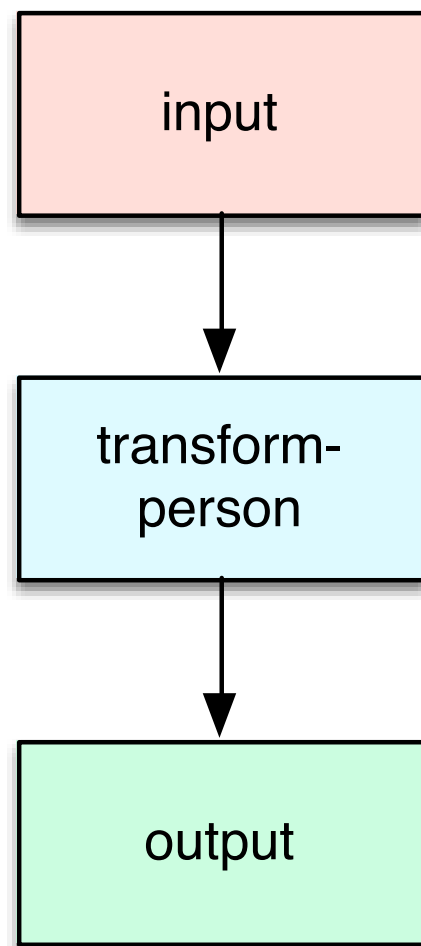
Positional Representation

- **Workflow:** describe the order of data flow through the program



Positional Representation

- **Workflow:** describe the order of data flow through the program

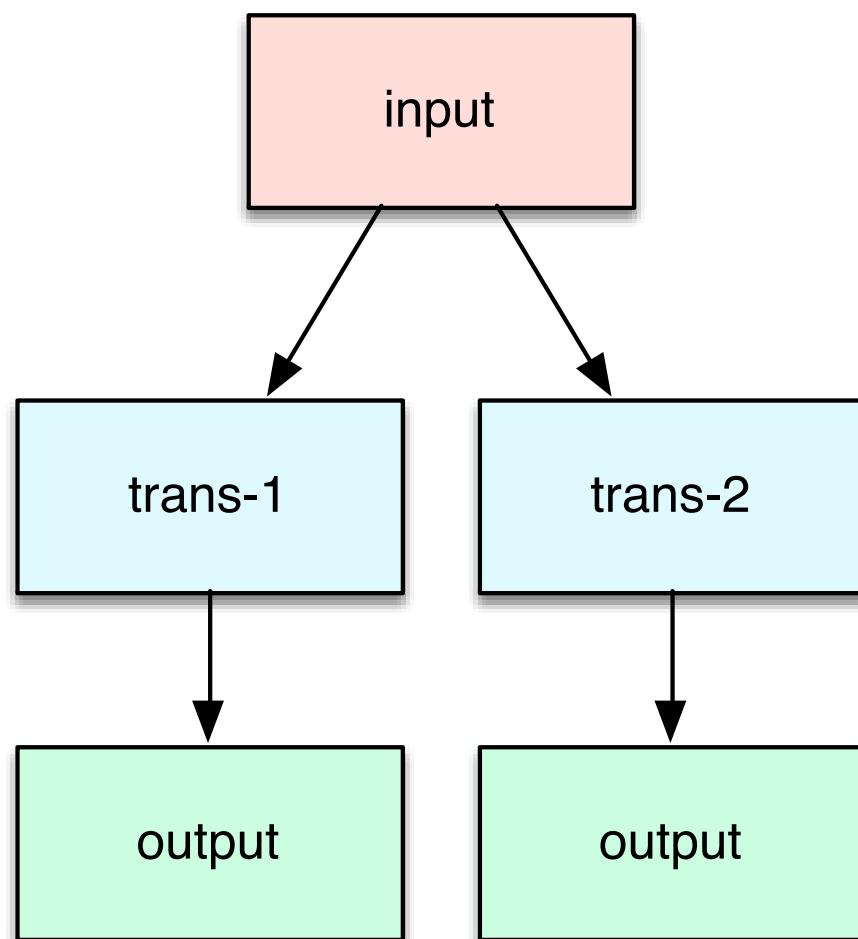


```
{:input {:transform-person :output}}
```



Positional Representation

- **Workflow:** describe the order of data flow through the program

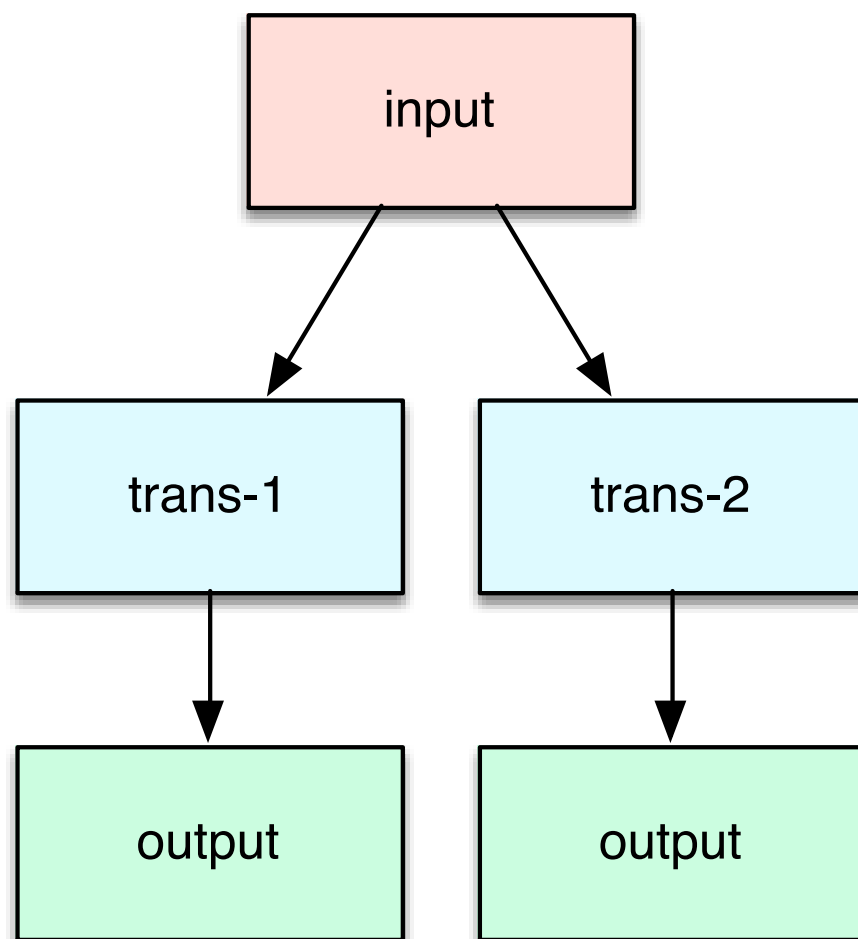


```
{:input {:trans-1 :output  
         :trans-2 :output}}
```



Positional Representation

- **Workflow:** describe the order of data flow through the program



```
{:input {:trans-1 :output  
         :trans-2 :output}}
```

language agnostic

location agnostic

temporally agnostic

machine tolerant



Execution Context

- **Catalog:** configure workflow tasks to implementation



workflow



```
{:in {:transform-person :out}}
```

catalog



```
[{:onyx/name :in
  :onyx/ident :mem/read-segments
  :onyx/type :input
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/bootstrap? true
  :onyx/batch-size 1024}

{:onyx/name :transform-person
  :onyx/fn :my.ns/transform-person
  :onyx/type :transformer
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}

{:onyx/name :out
  :onyx/ident :mem/write-segments
  :onyx/type :output
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}]
```

workflow



```
{:in {:transform-person :out}}
```

catalog



workflow



```
[{:onix/name :in  
  :onix/ident :mem/read-segments  
  :onix/type :input  
  :onix/medium :memory  
  :onix/consumption :concurrent  
  :onix/bootstrap? true  
  :onix/batch-size 1024}  
  
{:onix/name :transform-person  
  :onix/fn :my.ns/transform-person  
  :onix/type :transformer  
  :onix/consumption :concurrent  
  :onix/batch-size 1024}  
  
{:onix/name :out  
  :onix/ident :mem/write-segments  
  :onix/type :output  
  :onix/medium :memory  
  :onix/consumption :concurrent  
  :onix/batch-size 1024}]
```

`{:in {:transform-person :out}}`

catalog



```
[{:onix/name :in
  :onix/ident :mem/read-segments
  :onix/type :input
  :onix/medium :memory
  :onix/consumption :concurrent
  :onix/bootstrap? true
  :onix/batch-size 1024}


{:onix/name :transform-person
  :onix/fn :my.ns/transform-person
  :onix/type :transformer
  :onix/consumption :concurrent
  :onix/batch-size 1024}

{:onix/name :out
  :onix/ident :mem/write-segments
  :onix/type :output
  :onix/medium :memory
  :onix/consumption :concurrent
  :onix/batch-size 1024}]
```

workflow



```
{:in {:transform-person :out}}
```




catalog



```
[{:onyx/name :in
  :onyx/ident :mem/read-segments
  :onyx/type :input
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/bootstrap? true
  :onyx/batch-size 1024}

{:onyx/name :transform-person
  :onyx/fn :my.ns/transform-person
  :onyx/type :transformer
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}

{:onyx/name :out
  :onyx/ident :mem/write-segments
  :onyx/type :output
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}]
```



workflow



```
{:in {:transform-person :out}}
```

```
(defn transform-person [segment]
  (-> segment
    (update-in [:name] upper-case)
    (update-in [:age] inc)))
```

catalog



```
[{:onx/name :in
  :onx/ident :mem/read-segments
  :onx/type :input
  :onx/medium :memory
  :onx/consumption :concurrent
  :onx/bootstrap? true
  :onx/batch-size 1024}

{:onx/name :transform-person
  :onx/fn :my.ns/transform-person
  :onx/type :transformer
  :onx/consumption :concurrent
  :onx/batch-size 1024}

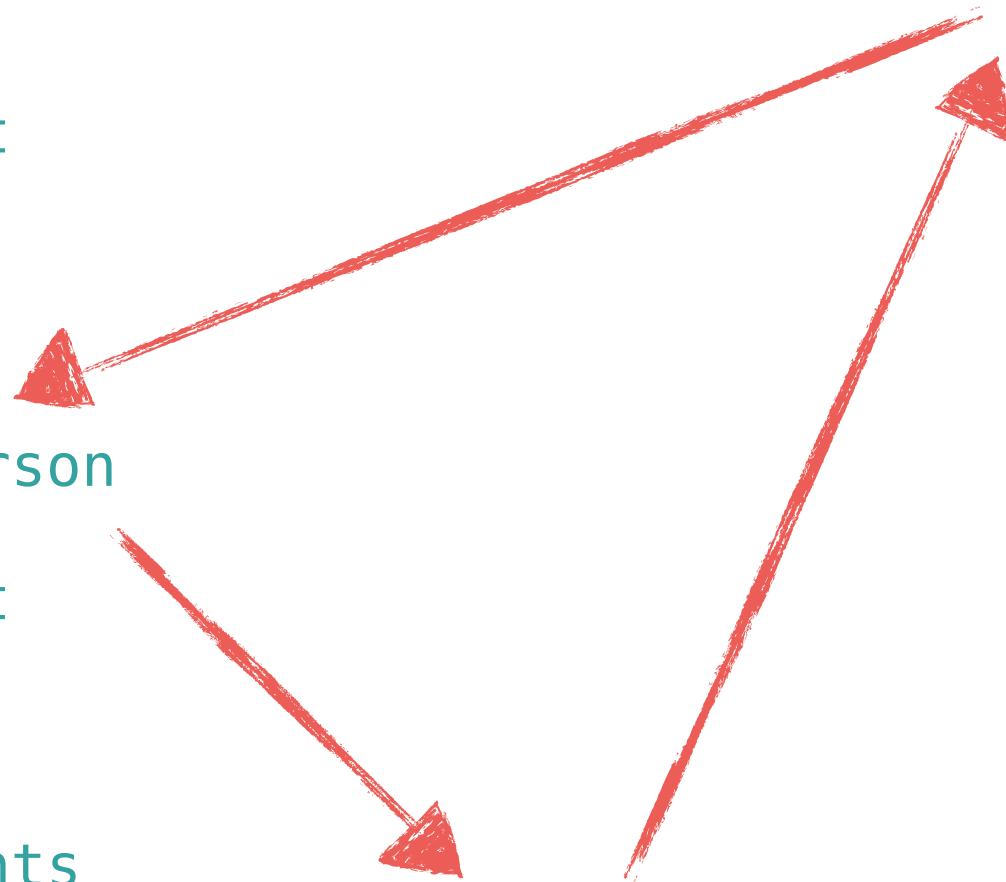
{:onx/name :out
  :onx/ident :mem/write-segments
  :onx/type :output
  :onx/medium :memory
  :onx/consumption :concurrent
  :onx/batch-size 1024}]
```

workflow



```
{:in {:transform-person :out}}
```

```
(defn transform-person [segment]
  (-> segment
    (update-in [:name] upper-case)
    (update-in [:age] inc)))
```



Extensibility Model

- Extensible to
 - Override things I don't like
 - Augment things I do like
- Reach to new storage mediums
 - With little to no code change in the application
- **Plugins:** adapt through the catalog



```
[{:onyx/name :in
  :onyx/ident :mem/read-segments
  :onyx/type :input
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/bootstrap? true
  :onyx/batch-size 1024}

{:onyx/name :transform-person
  :onyx/fn :my.ns/transform-person
  :onyx/type :transformer
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}

{:onyx/name :out
  :onyx/ident :mem/write-segments
  :onyx/type :output
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}]
```

```
[{:onyx/name :in
  :onyx/ident :mem/read-segments
  :onyx/type :input
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/bootstrap? true
  :onyx/batch-size 1024}

{:onyx/name :transform-person
  :onyx/fn :my.ns/transform-person
  :onyx/type :transformer
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}

{:onyx/name :out
  :onyx/ident :mem/write-segments
  :onyx/type :output
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}]
```

```
{:in {:transform-person :out}}
```

```
[{:onyx/name :in
  :onyx/ident :mem/read-segments
  :onyx/type :input
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/bootstrap? true
  :onyx/batch-size 1024}

{:onyx/name :transform-person
  :onyx/fn :my.ns/transform-person
  :onyx/type :transformer
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}

{:onyx/name :out
  :onyx/ident :mem/write-segments
  :onyx/type :output
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}]
```

```
{:in {:transform-person :out}}
```

```
[{:onyx/name :in
  :onyx/ident :mem/read-segments
  :onyx/type :input
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/bootstrap? true
  :onyx/batch-size 1024}

{:onyx/name :transform-person
  :onyx/fn :my.ns/transform-person
  :onyx/type :transformer
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}

{:onyx/name :out
  :onyx/ident :mem/write-segments
  :onyx/type :output
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}]
```

```
{:onyx/name :in
  :onyx/ident :hornetq/read-segments
  :onyx/type :input
  :onyx/medium :hornetq
  :onyx/consumption :concurrent
  :kafka/topic "my-topic"
  :kafka/zookeeper "127.0.0.1:2181"
  :kafka/group-id "onyx-consumer"
  :kafka/offset-reset "smallest"
  :onyx/batch-size 1024}
```

```
{:in {:transform-person :out}}
```

```
[{:onyx/name :in
  :onyx/ident :mem/read-segments
  :onyx/type :input
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/bootstrap? true
  :onyx/batch-size 1024}

{:onyx/name :transform-person
  :onyx/fn :my.ns/transform-person
  :onyx/type :transformer
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}

{:onyx/name :out
  :onyx/ident :mem/write-segments
  :onyx/type :output
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}]
```



```
{:onyx/name :in
  :onyx/ident :hornetq/read-segments
  :onyx/type :input
  :onyx/medium :hornetq
  :onyx/consumption :concurrent
  :kafka/topic "my-topic"
  :kafka/zookeeper "127.0.0.1:2181"
  :kafka/group-id "onyx-consumer"
  :kafka/offset-reset "smallest"
  :onyx/batch-size 1024}
```

```
{:in {:transform-person :out}}
```

```
[{:onyx/name :in
  :onyx/ident :mem/read-segments
  :onyx/type :input
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/bootstrap? true
  :onyx/batch-size 1024}

{:onyx/name :transform-person
  :onyx/fn :my.ns/transform-person
  :onyx/type :transformer
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}

{:onyx/name :out
  :onyx/ident :mem/write-segments
  :onyx/type :output
  :onyx/medium :memory
  :onyx/consumption :concurrent
  :onyx/batch-size 1024}]
```

```
{:onyx/name :in
  :onyx/ident :hornetq/read-segments
  :onyx/type :input
  :onyx/medium :hornetq
  :onyx/consumption :concurrent
  :kafka/topic "my-topic"
  :kafka/zookeeper "127.0.0.1:2181"
  :kafka/group-id "onyx-consumer"
  :kafka/offset-reset "smallest"
  :onyx/batch-size 1024}

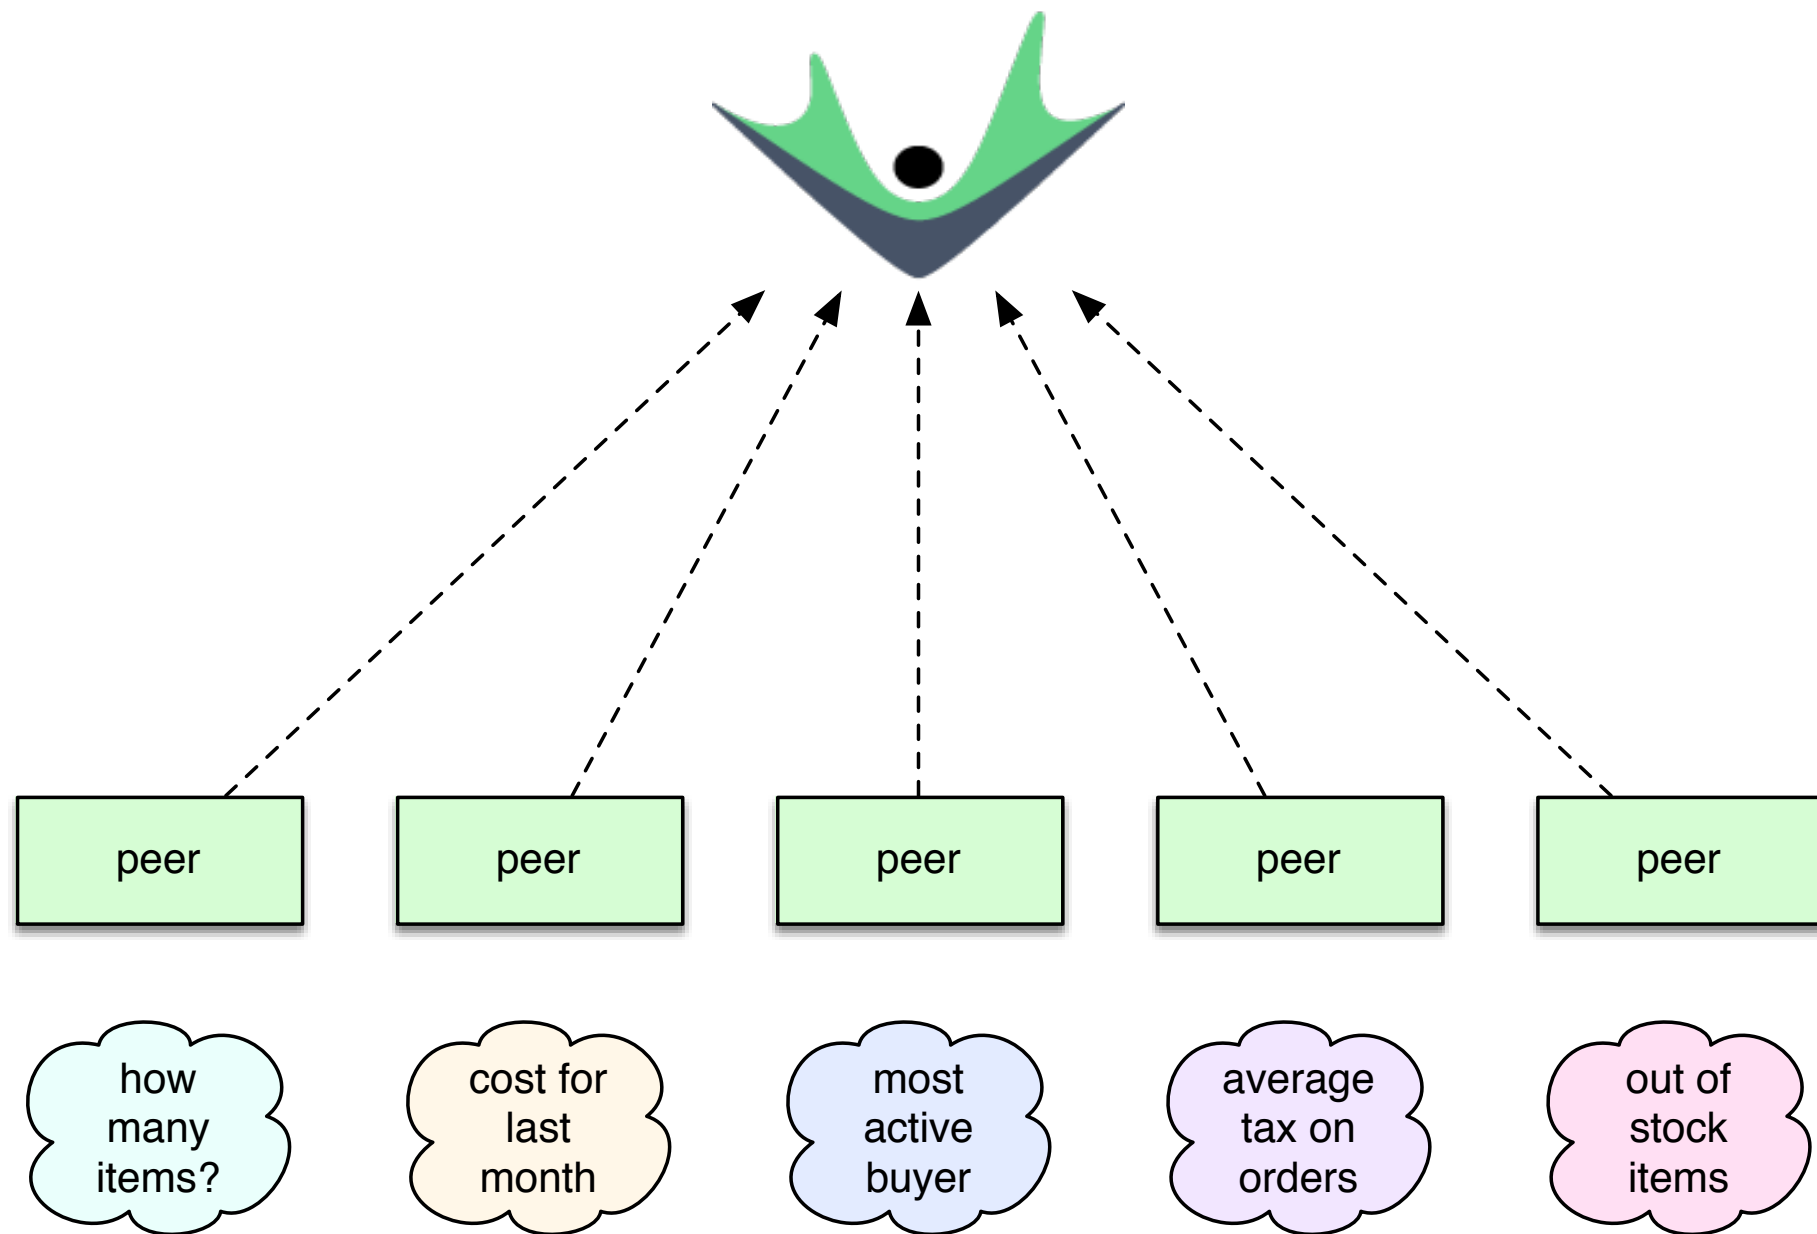
{:in {:transform-person :out}}
```



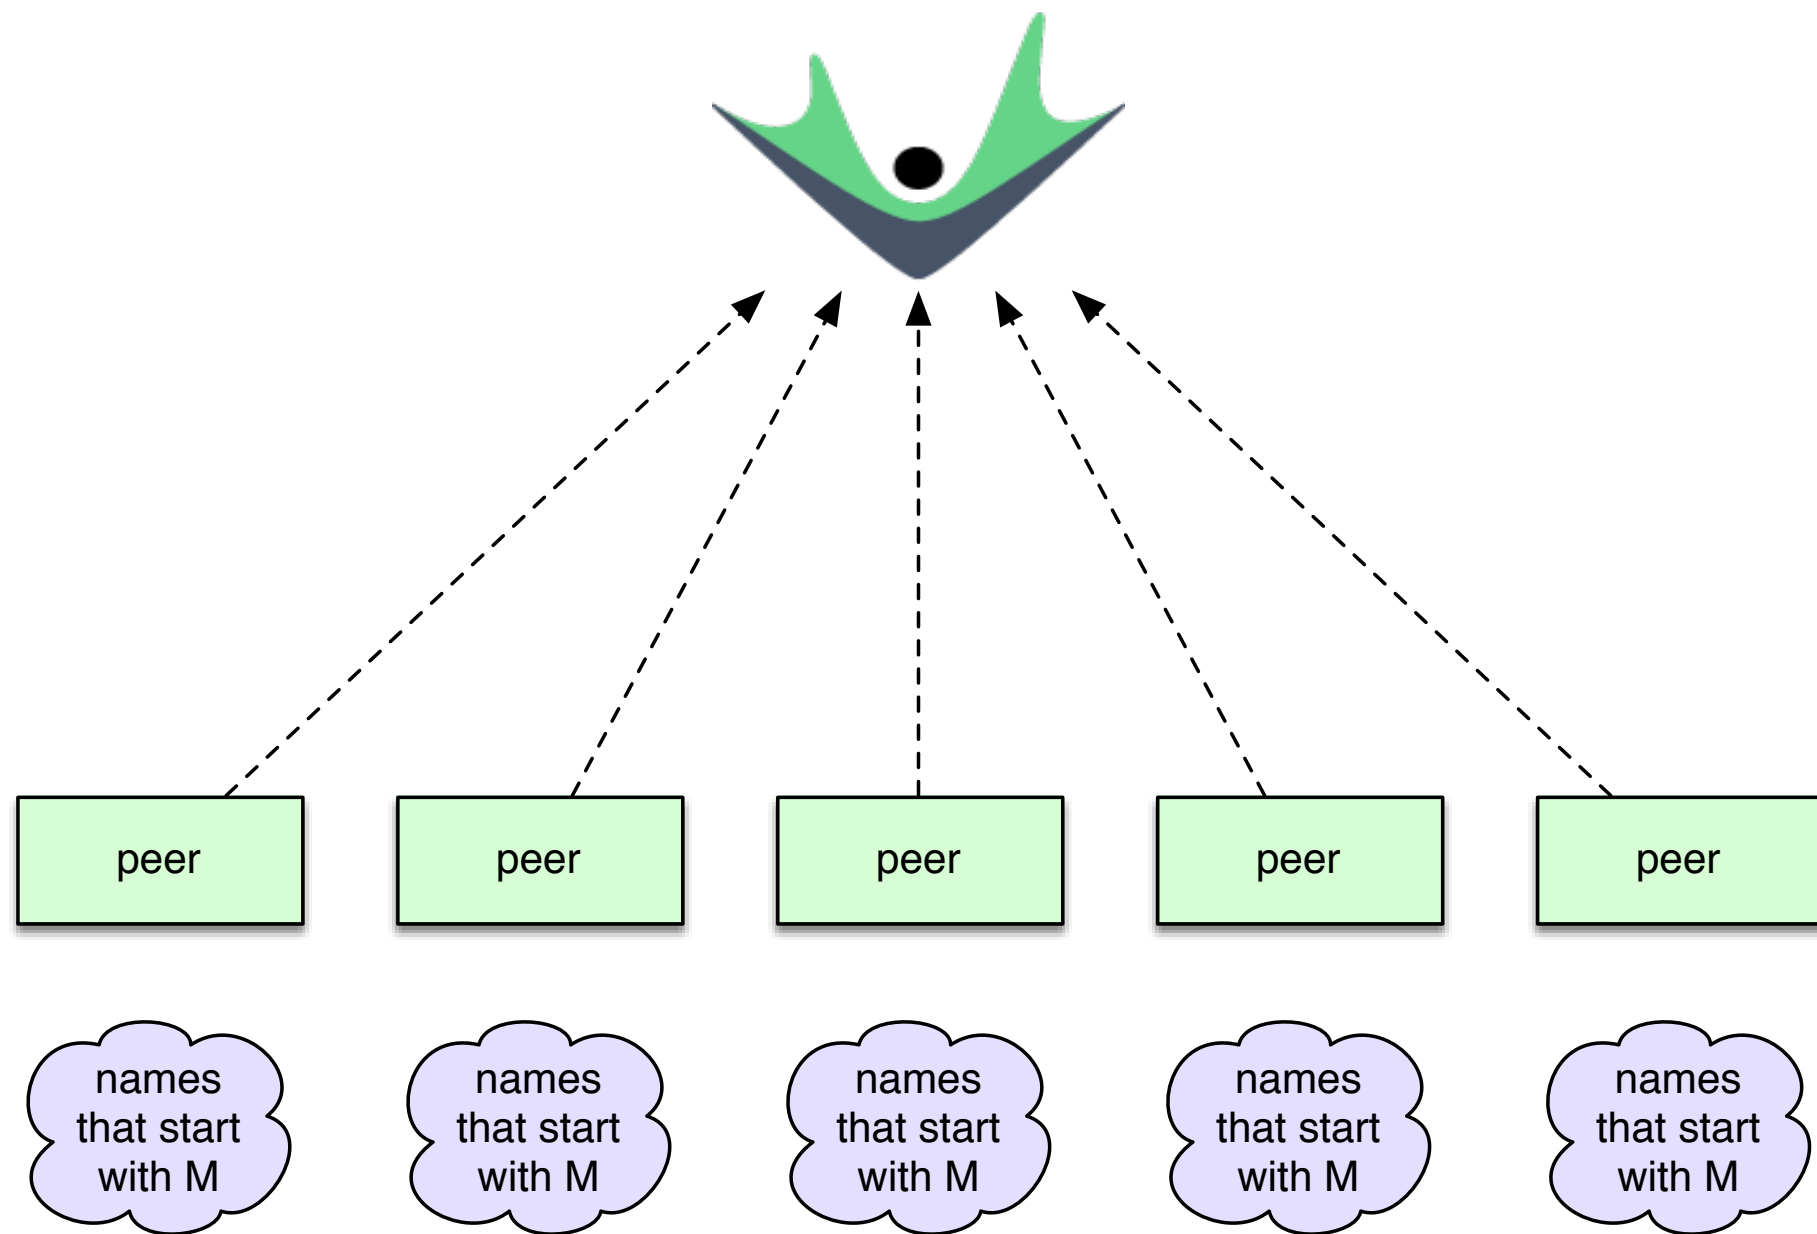
Runtime Discovery



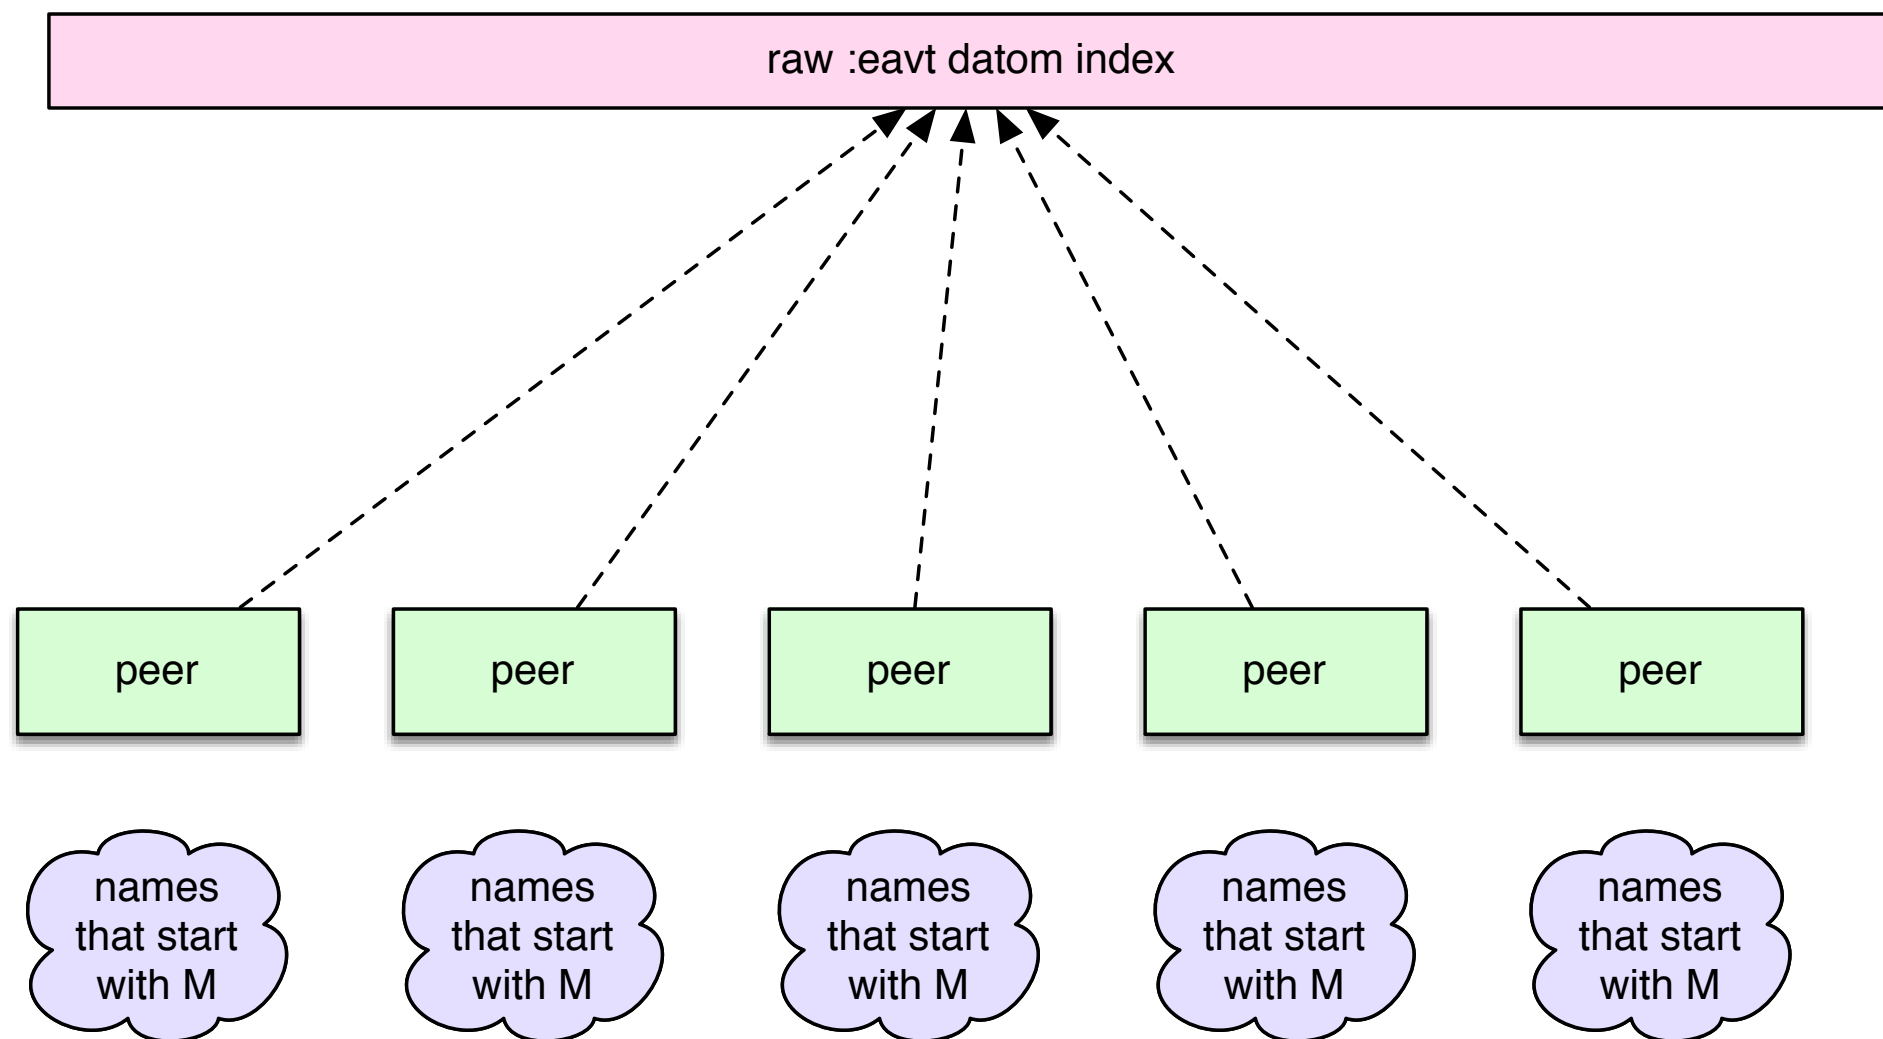
Runtime Discovery



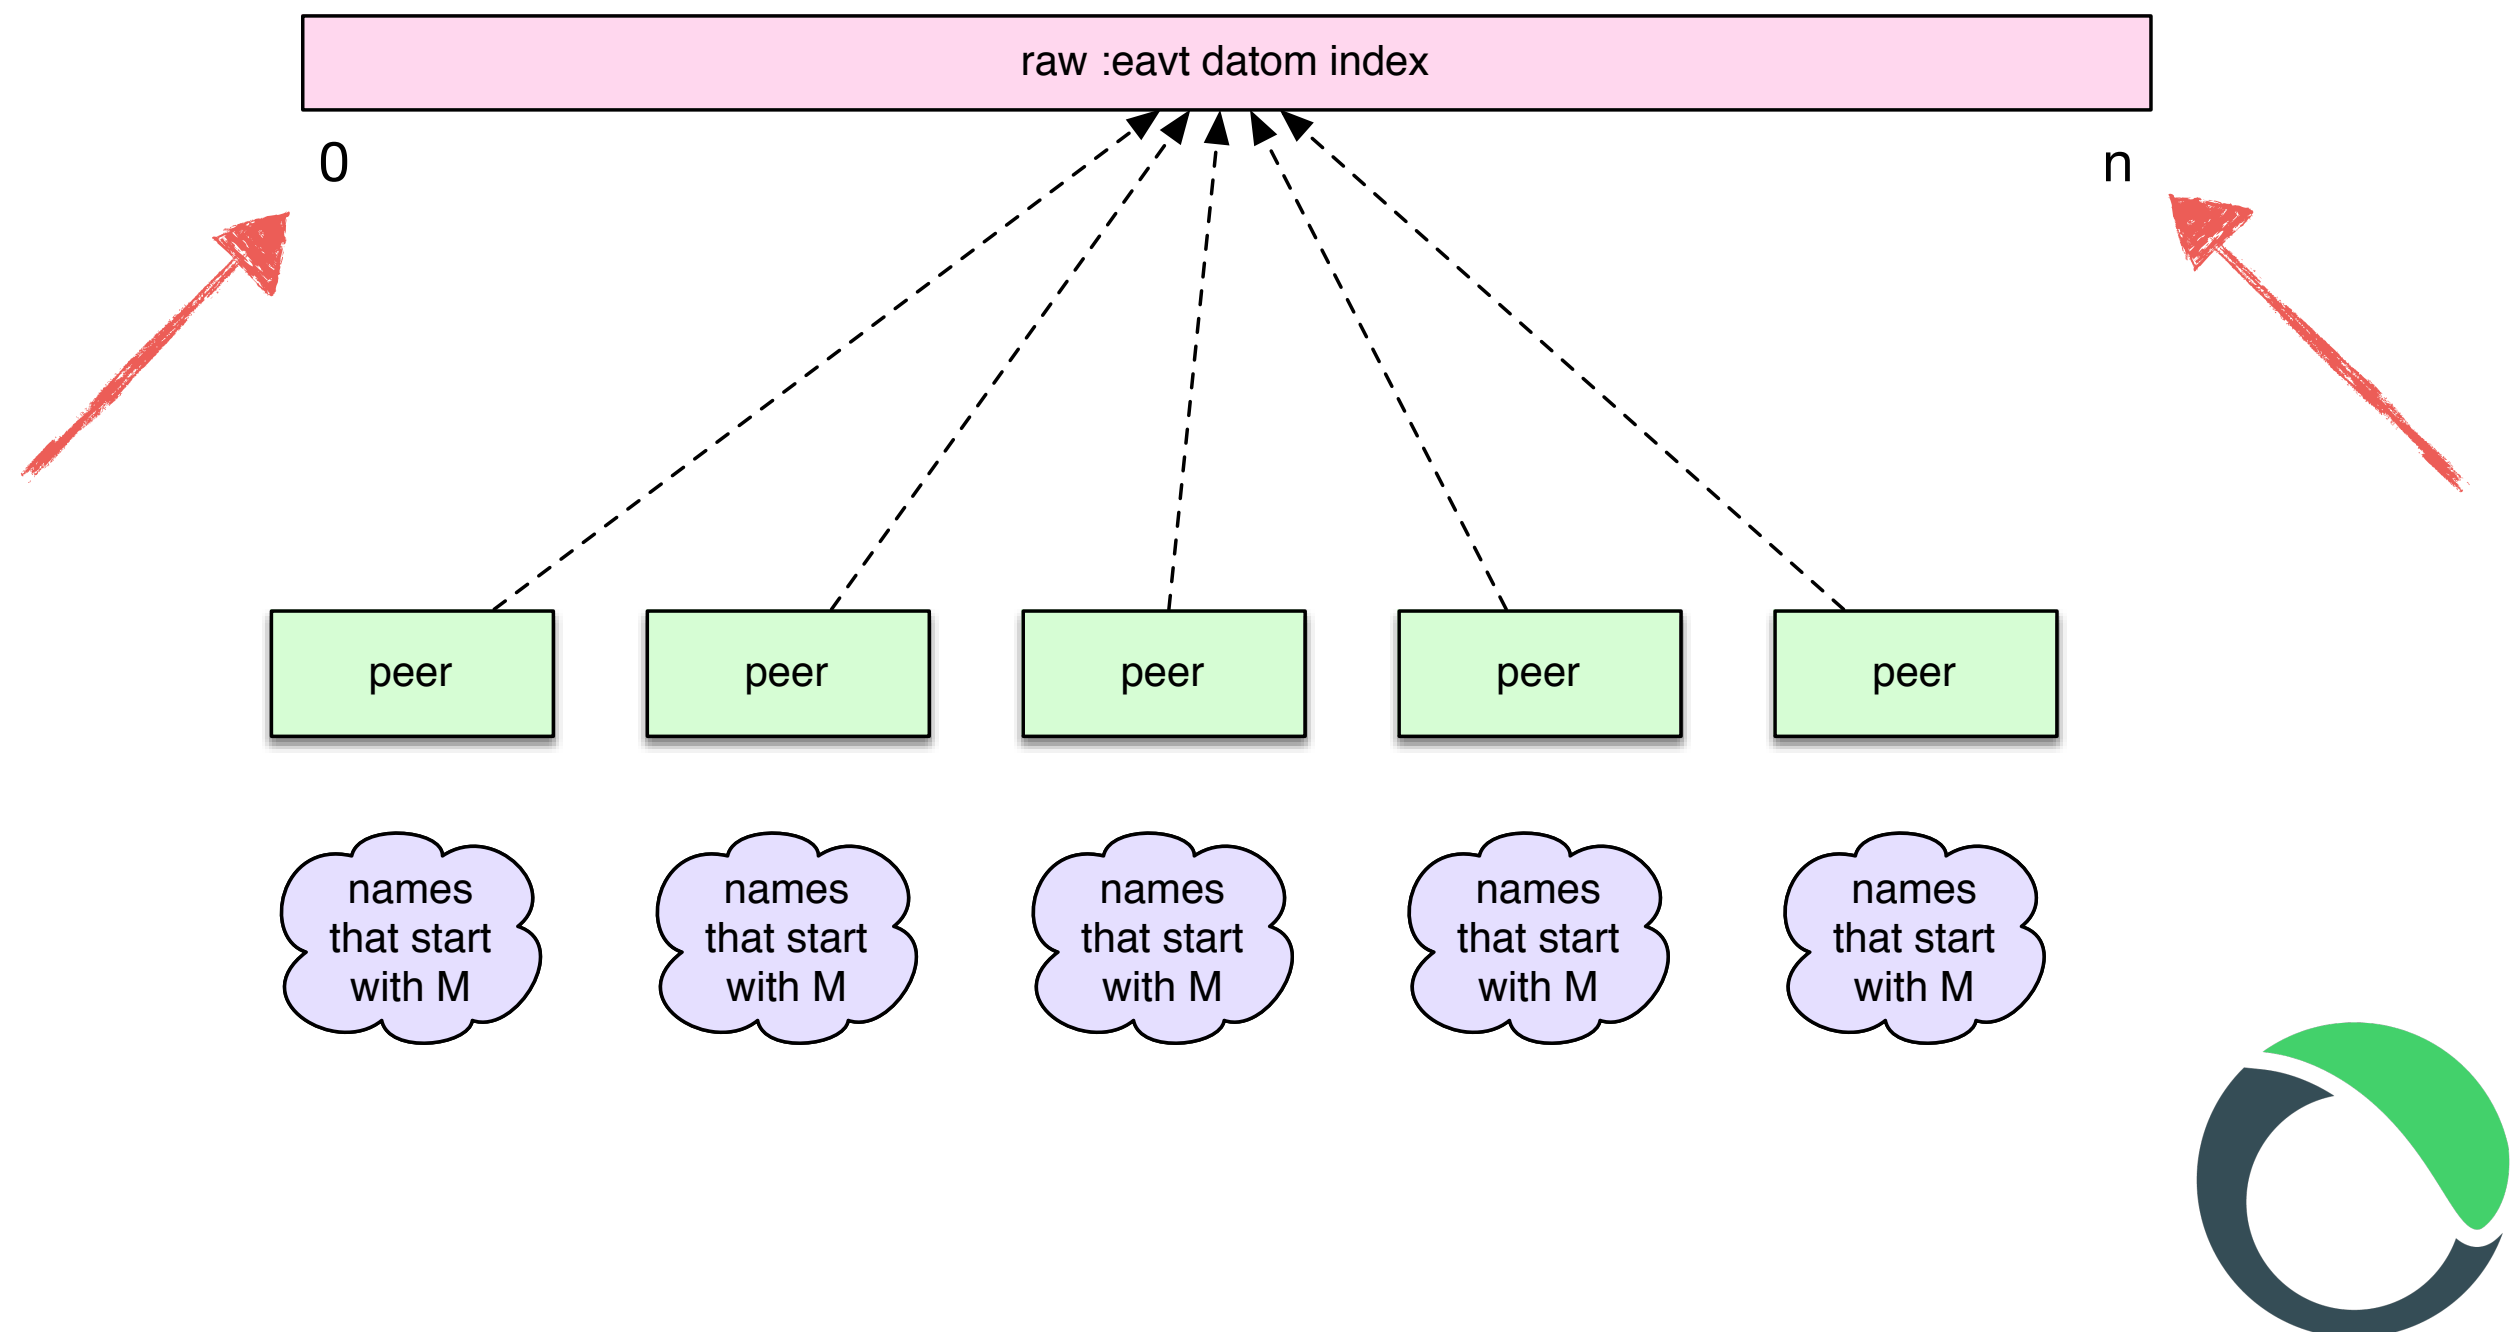
Runtime Discovery



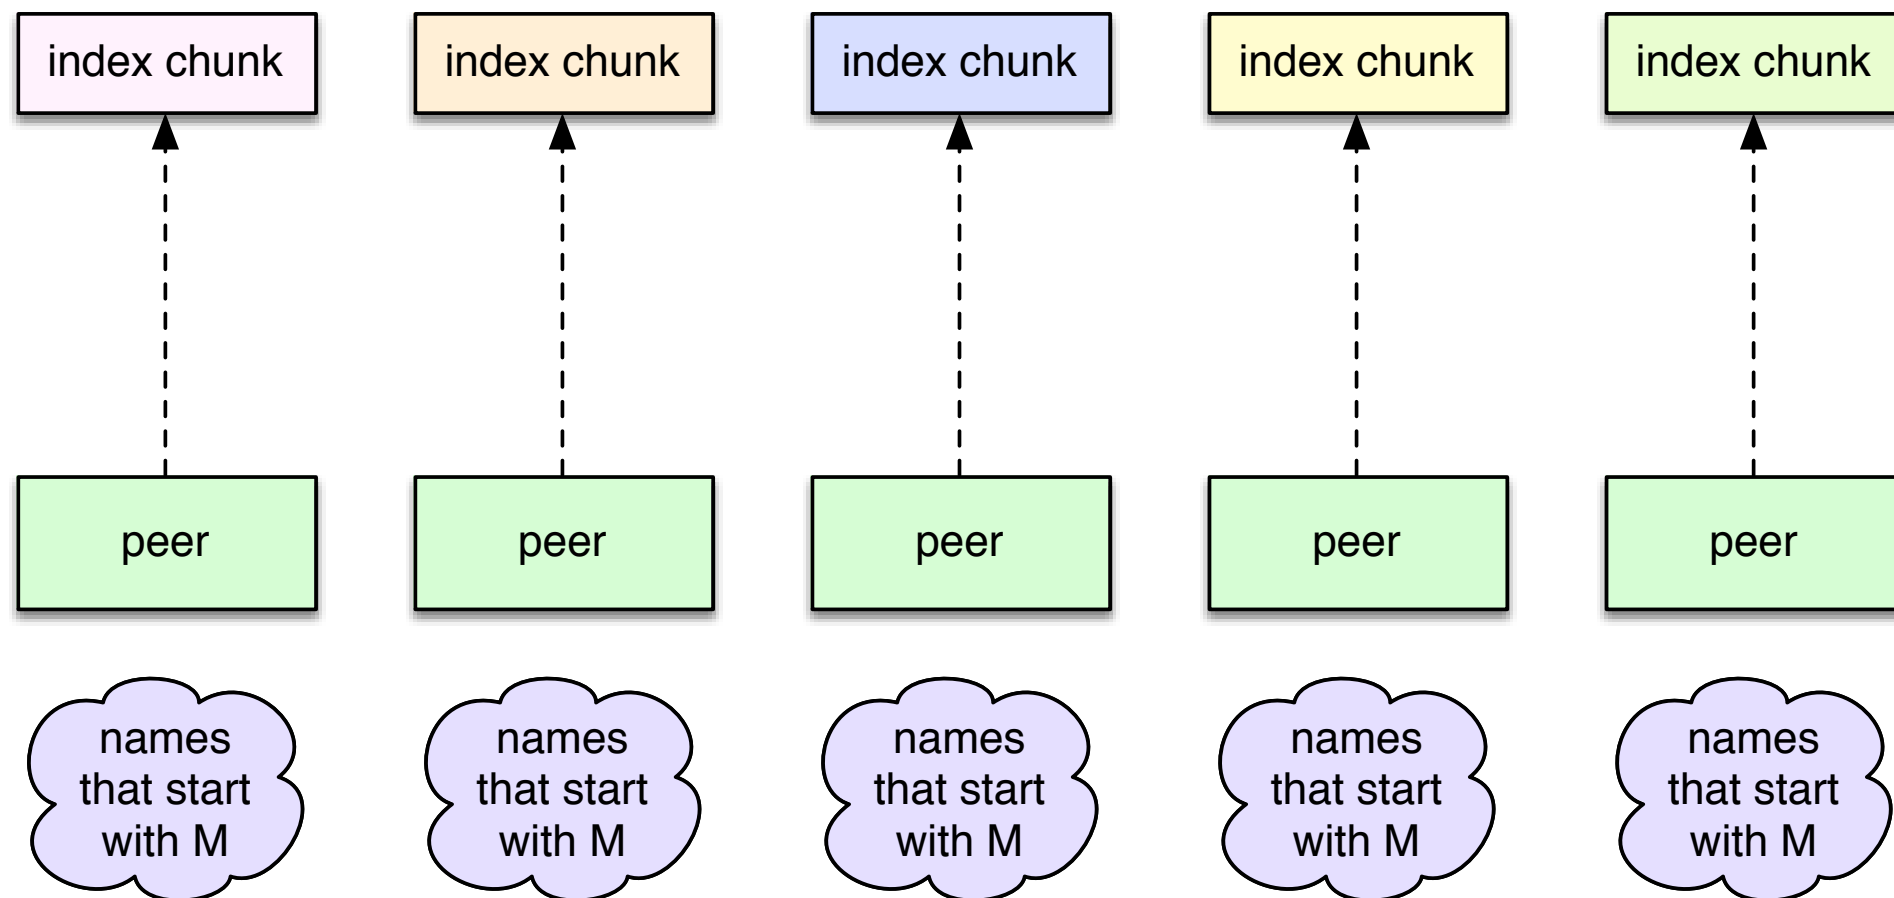
Runtime Discovery



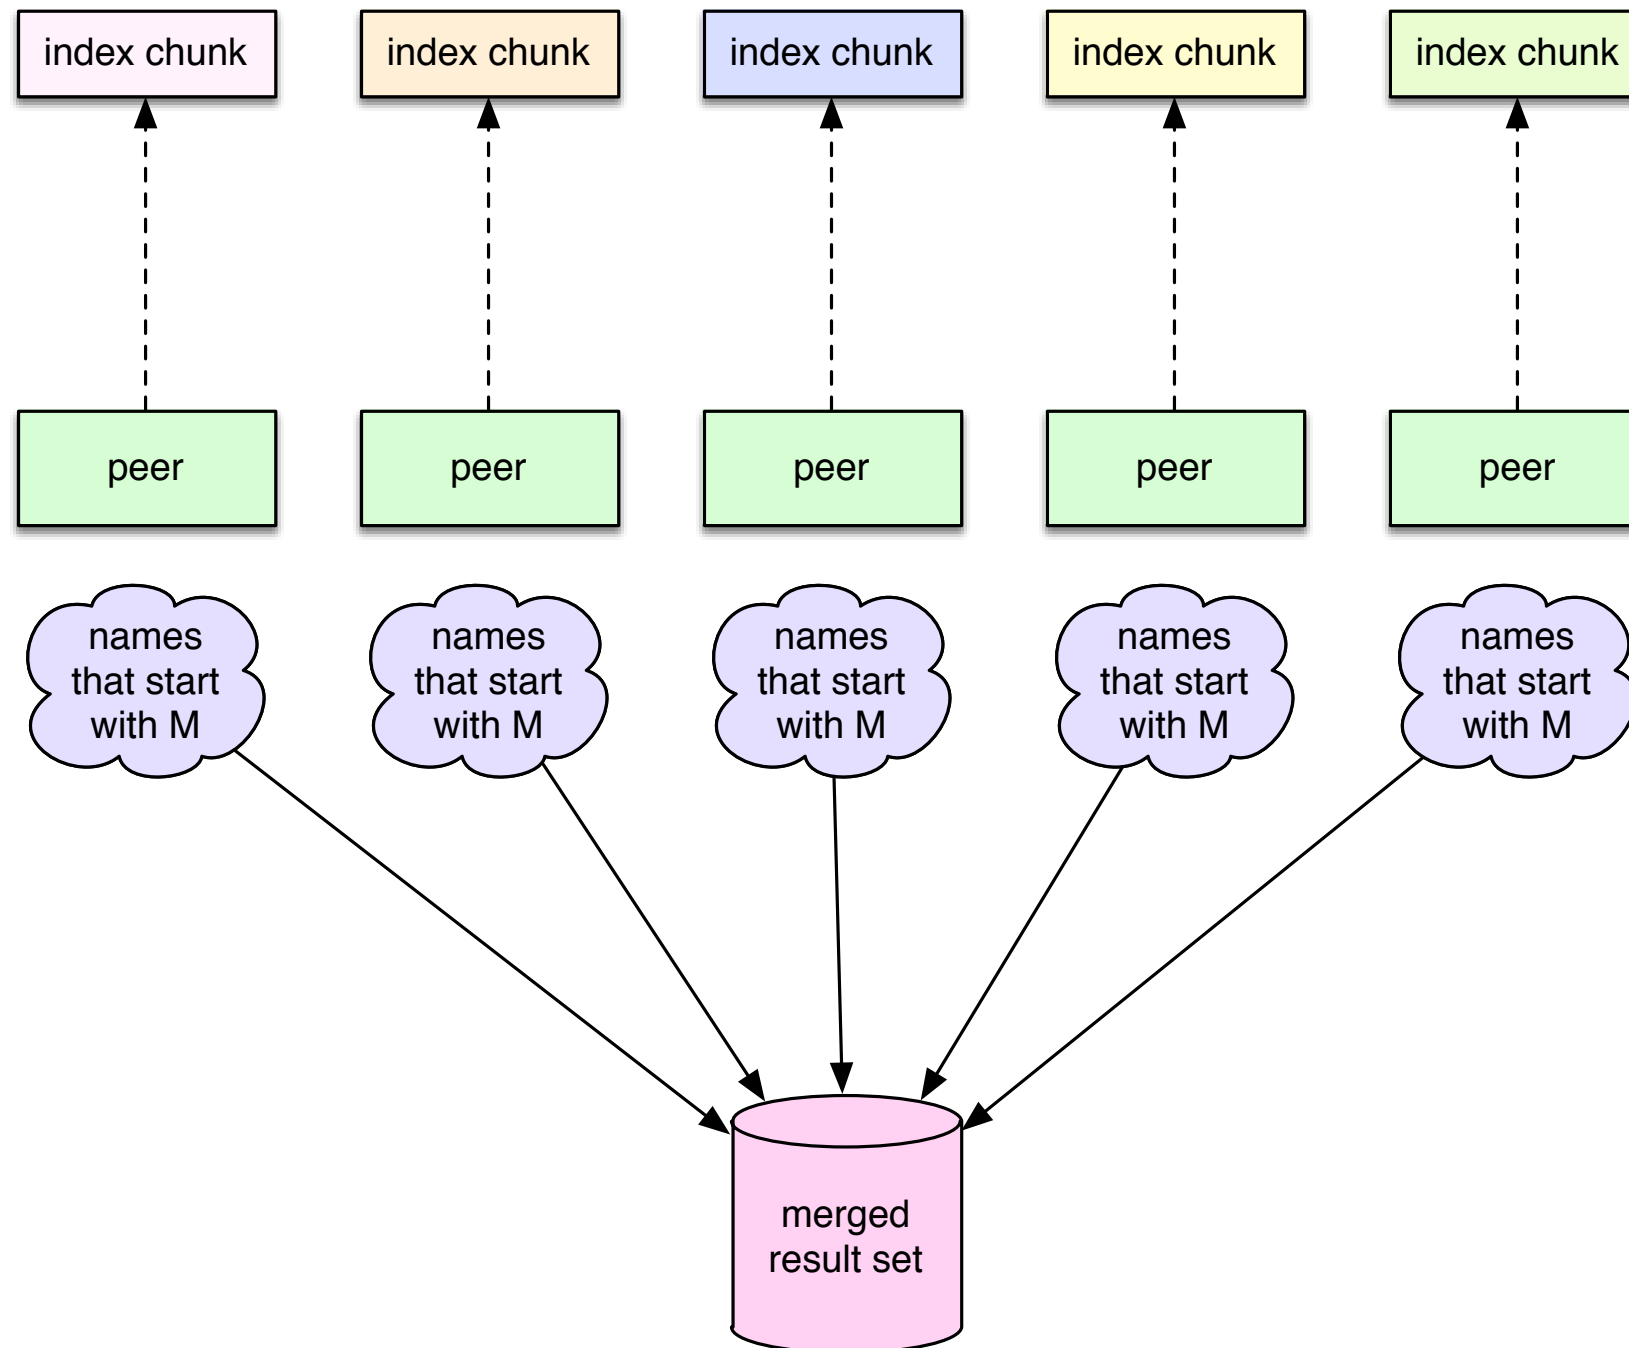
Runtime Discovery



Runtime Discovery



Runtime Discovery



Runtime Discovery

```
{:partition-datoms {:read-datoms {:query :persist}}}
```



Runtime Discovery

```
{:partition-datoms {:read-datoms {:query :persist}}}
```

```
[{:onyx/name :partition-datoms  
 :onyx/ident :datomic/partition-datoms  
 :onyx/type :input  
 :onyx/medium :datomic  
 :onyx/consumption :sequential  
 :onyx/bootstrap? true  
 :datomic/uri db-uri  
 :datomic/t t  
 :datomic/datoms-per-segment datom-size  
 :datomic/partition :com.mdrogalis/people  
 :onyx/batch-size batch-size}]
```


```
{:onyx/name :read-datoms  
 :onyx/ident :datomic/read-datoms  
 :onyx/fn :onyx.plugin.datomic/read-datoms  
 :onyx/type :transformer  
 :onyx/consumption :concurrent  
 :onyx/batch-size batch-size  
 :datomic/uri db-uri  
 :datomic/t t}
```



Runtime Discovery

```
{:partition-datoms {:read-datoms {:query :persist}}}
```

```
[{:onyx/name :partition-datoms  
 :onyx/ident :datomic/partition-datoms  
 :onyx/type :input  
 :onyx/medium :datomic  
 :onyx/consumption :sequential  
 :onyx/bootstrap? true  
 :datomic/uri db-uri  
 :datomic/t t  
 :datomic/datoms-per-segment datom-size  
 :datomic/partition :com.mdrogalis/people  
 :onyx/batch-size batch-size}
```




```
{:onyx/name :read-datoms  
 :onyx/ident :datomic/read-datoms  
 :onyx/fn :onyx.plugin.datomic/read-datoms  
 :onyx/type :transformer  
 :onyx/consumption :concurrent  
 :onyx/batch-size batch-size  
 :datomic/uri db-uri  
 :datomic/t t}
```



Runtime Discovery

```
{:partition-datoms {:read-datoms {:query :persist}}}
```

```
[{:onyx/name :partition-datoms  
 :onyx/ident :datomic/partition-datoms  
 :onyx/type :input  
 :onyx/medium :datomic  
 :onyx/consumption :sequential  
 :onyx/bootstrap? true  
 :datomic/uri db-uri  
 :datomic/t t  
 :datomic/datoms-per-segment datom-size  
 :datomic/partition :com.mdrogalis/people  
 :onyx/batch-size batch-size}]
```



```
{:onyx/name :read-datoms  
 :onyx/ident :datomic/read-datoms  
 :onyx/fn :onyx.plugin.datomic/read-datoms  
 :onyx/type :transformer  
 :onyx/consumption :concurrent  
 :onyx/batch-size batch-size  
 :datomic/uri db-uri  
 :datomic/t t}
```



Runtime Discovery

```
{:partition-datoms {:read-datoms {:query :persist}}}
```

```
[{:onyx/name :partition-datoms  
 :onyx/ident :datomic/partition-datoms  
 :onyx/type :input  
 :onyx/medium :datomic  
 :onyx/consumption :sequential  
 :onyx/bootstrap? true  
 :datomic/uri db-uri  
 :datomic/t t  
 :datomic/datoms-per-segment datom-size  
 :datomic/partition :com.mdrogalis/people  
 :onyx/batch-size batch-size}]
```

```
{:onyx/name :read-datoms  
 :onyx/ident :datomic/read-datoms  
 :onyx/fn :onyx.plugin.datomic/read-datoms  
 :onyx/type :transformer  
 :onyx/consumption :concurrent  
 :onyx/batch-size batch-size  
 :datomic/uri db-uri  
 :datomic/t t}
```



Runtime Discovery

```
{:partition-datoms {:read-datoms {:query :persist}}}
```

```
[{:onx/name :partition-datoms  
 :onx/ident :datomic/partition-datoms  
 :onx/type :input  
 :onx/medium :datomic  
 :onx/consumption :sequential  
 :onx/bootstrap? true  
 :datomic/uri db-uri  
 :datomic/t t  
 :datomic/datoms-per-segment datom-size  
 :datomic/partition :com.mdrogalis/people  
 :onx/batch-size batch-size}
```

```
{:onx/name :read-datoms  
 :onx/ident :datomic/read-datoms  
 :onx/fn :onx.plugin.datomic/read-datoms  
 :onx/type :transformer  
 :onx/consumption :concurrent  
 :onx/batch-size batch-size  
 :datomic/uri db-uri  
 :datomic/t t}
```

```
(defn my-query [{:keys [datoms]}  
  {:result (d/q query datoms)})
```



Task State API

- Your program *needs* side-effects
 - database connection, file handles, ...
- def/defonce - uh oh
- **Lifecycles**: managed set-up/teardown of state



Task State API

`(defmulti start-lifecycle?)`

`(defmulti inject-lifecycle-resources)`

`(defmulti inject-temporal-resources)`

`(defmulti close-temporal-resources)`

`(defmulti close-lifecycle-resources)`



Task State API

(defmulti start-lifecycle?)

(defmulti inject-lifecycle-resources)

(defmulti inject-temporal-resources)

(defmulti close-temporal-resources)

(defmulti close-lifecycle-resources)



Task State API

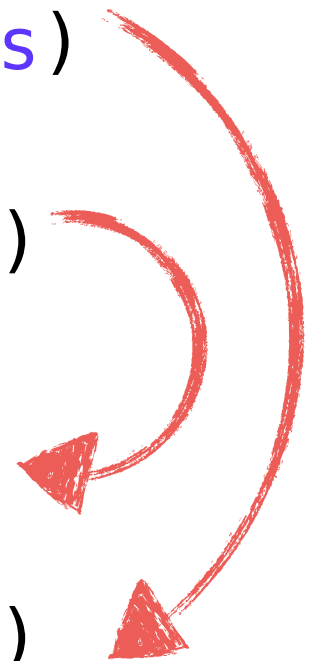
(defmulti start-lifecycle?)

(defmulti inject-lifecycle-resources)

(defmulti inject-temporal-resources)

(defmulti close-temporal-resources)

(defmulti close-lifecycle-resources)




Task State API

```
{:onyx/name :my-task  
 :onyx/ident :strangeloop/task  
 :onyx/fn :ns.strangeloop/task  
 :onyx/type :transformer  
 :onyx/consumption :concurrent  
 :my/param "42"  
 :onyx/batch-size 1024}
```



Task State API

```
{:onyx/name :my-task  
 :onyx/ident :strangeloop/task  
 :onyx/fn :ns.strangeloop/task  
 :onyx/type :transformer  
 :onyx/consumption :concurrent  
 :my/param "42"  
 :onyx/batch-size 1024}
```



Task State API

```
{:onyx/name :my-task  
 :onyx/ident :strangeloop/task  
 :onyx/fn :ns.strangeloop/task  
 :onyx/type :transformer  
 :onyx/consumption :concurrent  
 :my/param "42"  
 :onyx/batch-size 1024}
```

```
(defmethod l-ext/inject-lifecycle-resources  
 :strangeloop/task  
 [_ {:keys [onyx.core/task-map]}]  
 {:my-thing (create-stateful-thing (:my/param task-map))})
```



Task State API

```
{:onyx/name :my-task  
 :onyx/ident :strangeloop/task  
 :onyx/fn :ns.strangeloop/task  
 :onyx/type :transformer  
 :onyx/consumption :concurrent  
 :my/param "42"  
 :onyx/batch-size 1024}
```

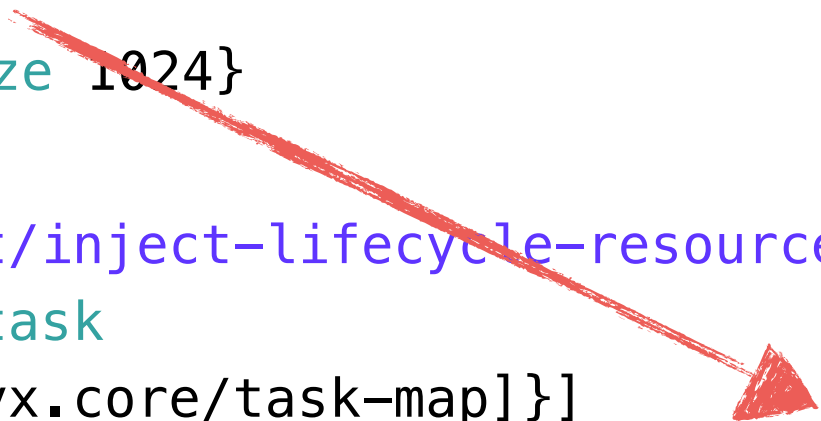
```
(defmethod l-ext/inject-lifecycle-resources  
  :strangeloop/task  
  [_ {:keys [onyx.core/task-map]}]  
  {:my-thing (create-stateful-thing (:my/param task-map))})
```



Task State API

```
{:onyx/name :my-task  
 :onyx/ident :strangeloop/task  
 :onyx/fn :ns.strangeloop/task  
 :onyx/type :transformer  
 :onyx/consumption :concurrent  
 :my/param "42"  
 :onyx/batch-size 1024}
```

```
(defmethod l-ext/inject-lifecycle-resources  
 :strangeloop/task  
 [_ {:keys [onyx.core/task-map]}]  
 {:my-thing (create-stateful-thing (:my/param task-map))})
```



Task State API

```
{:onyx/name :my-task  
 :onyx/ident :strangeloop/task  
 :onyx/fn :ns.strangeloop/task  
 :onyx/type :transformer  
 :onyx/consumption :concurrent  
 :my/param "42"  
 :onyx/batch-size 1024}
```

```
(defmethod l-ext/inject-lifecycle-resources  
 :strangeloop/task  
 [_ {:keys [onyx.core/task-map]}]  
 {:my-thing (create-stateful-thing (:my/param task-map))})
```



Task State API

```
{:onyx/name :my-task  
 :onyx/ident :strangeloop/task  
 :onyx/fn :ns.strangeloop/task  
 :onyx/type :transformer  
 :onyx/consumption :concurrent  
 :my/param "42"  
 :onyx/batch-size 1024}
```

```
(defmethod l-ext/inject-lifecycle-resources  
  :strangeloop/task  
  [_ {:keys [onyx.core/task-map]}]  
  {:my-thing (create-stateful-thing (:my/param task-map))})
```

```
(defmethod l-ext/close-temporal-resources  
  :strangeloop/task  
  [_ event]  
  (do-side-effects (:my-thing event))  
  {})
```

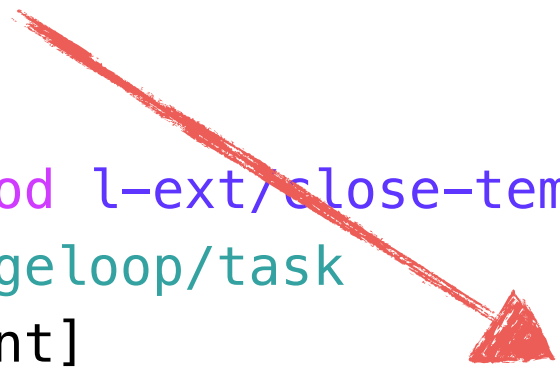


Task State API

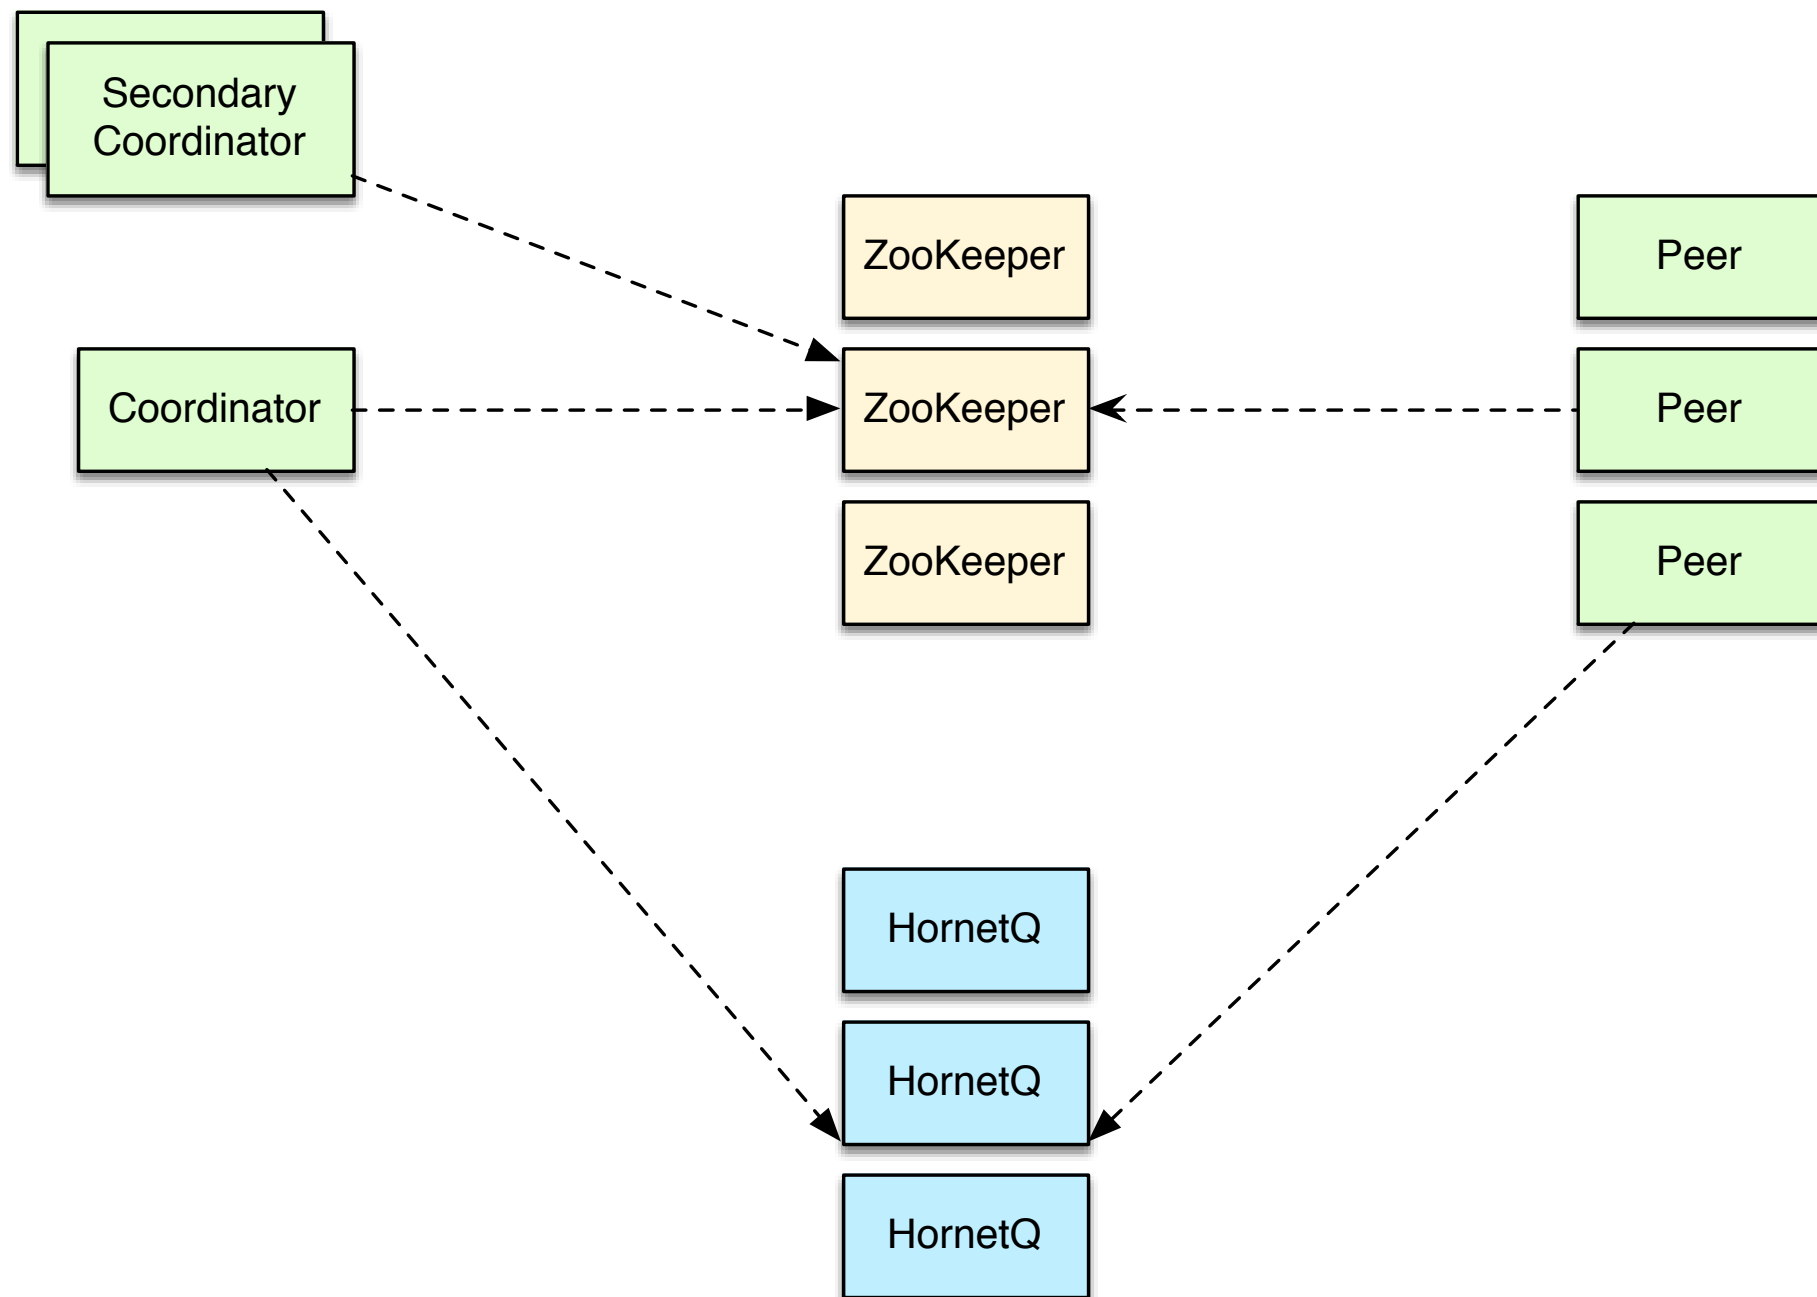
```
{:onyx/name :my-task  
 :onyx/ident :strangeloop/task  
 :onyx/fn :ns.strangeloop/task  
 :onyx/type :transformer  
 :onyx/consumption :concurrent  
 :my/param "42"  
 :onyx/batch-size 1024}
```

```
(defmethod l-ext/inject-lifecycle-resources  
 :strangeloop/task  
 [_ {:keys [onyx.core/task-map]}]  
 {:my-thing (create-stateful-thing (:my/param task-map))})
```

```
(defmethod l-ext/close-temporal-resources  
 :strangeloop/task  
 [_ event]  
 (do-side-effects (:my-thing event))  
 {})
```



Architecture Overview



Wait! That's Slow!

- Queue-based architectures criticized for speed
- Disk locality within data center irrelevant (Berkeley)
- 10+ gig switches now available in data centers



Demo!



Thanks

- CircleCI
- Infinite Cloud



Questions?



<https://github.com/MichaelDrogalis/onyx>

@MichaelDrogalis