

Elixir and the Internet of Things

Handling a Stampede



Doug Rohrer

CC by <https://www.flickr.com/people/t3rmin4t0r/>

Me

- Professional software developer for over 20 years
- Written in many languages on many operating systems
- Often called upon to understand and mitigate performance issues on projects
- Had never touched Elixir or Erlang (or much FP) before this project

The Problem

- 10s to 100s of thousands of connected devices
- Sometimes all connecting in a short period of time
- Need to process many messages in a short period of time (round-trip time is critical, as users are waiting for doors to unlock)
- Beyond functionally routing messages, the number one requirement for this application was to scale to very large (200k) numbers of simultaneous connections

Our Process

- Built an performance-testing tool (in Go)
- Built a potential solution (Ruby, Elixir)
- Test
- Profiled applications to find and fix bottlenecks, and retest
- Abandon the solution when there were no clear bottlenecks to fix but performance was inadequate

We Tried Ruby/EventMachine

- Pros
 - Existing solution
- Cons
 - Single-threaded event loop (callback hell)
 - High CPU overhead for AMQP
 - Unbalanced I/O - Would do all of “I” before any “O”

We Tried Ruby/Celluloid.io

- Pros
 - Agents are easier to understand than lots of callbacks and `next_tick` calls
- Cons
 - Uses Fibers, which are native threads on JRuby
 - Uses LOTS of memory (40GB for ~40k connections)

Why Try Erlang?

- Erlang/OTP was designed to handle this kind of problem
- `{ok, concurrency} = erlang:use()`
 - Concurrency is built into the runtime
 - Actor model is easy to implement
 - Known for handling the load levels we needed
- Clustering is (almost) code free
 - We need to load balance across machines
 - Erlang allows us to do this with minimal additional coding

Why Elixir and not Erlang?

- The client is predominantly a Ruby shop, and Elixir was a more comfortable choice for them as it has a more familiar syntax than Erlang
- Macros reduce “template code” bloat & potential for errors
- Can still use all Erlang libraries
- Compiles to Erlang bytecode and runs on the Erlang VM
- For all practical purposes, Elixir was \geq Erlang for us

Some Erlang/Elixir Basics



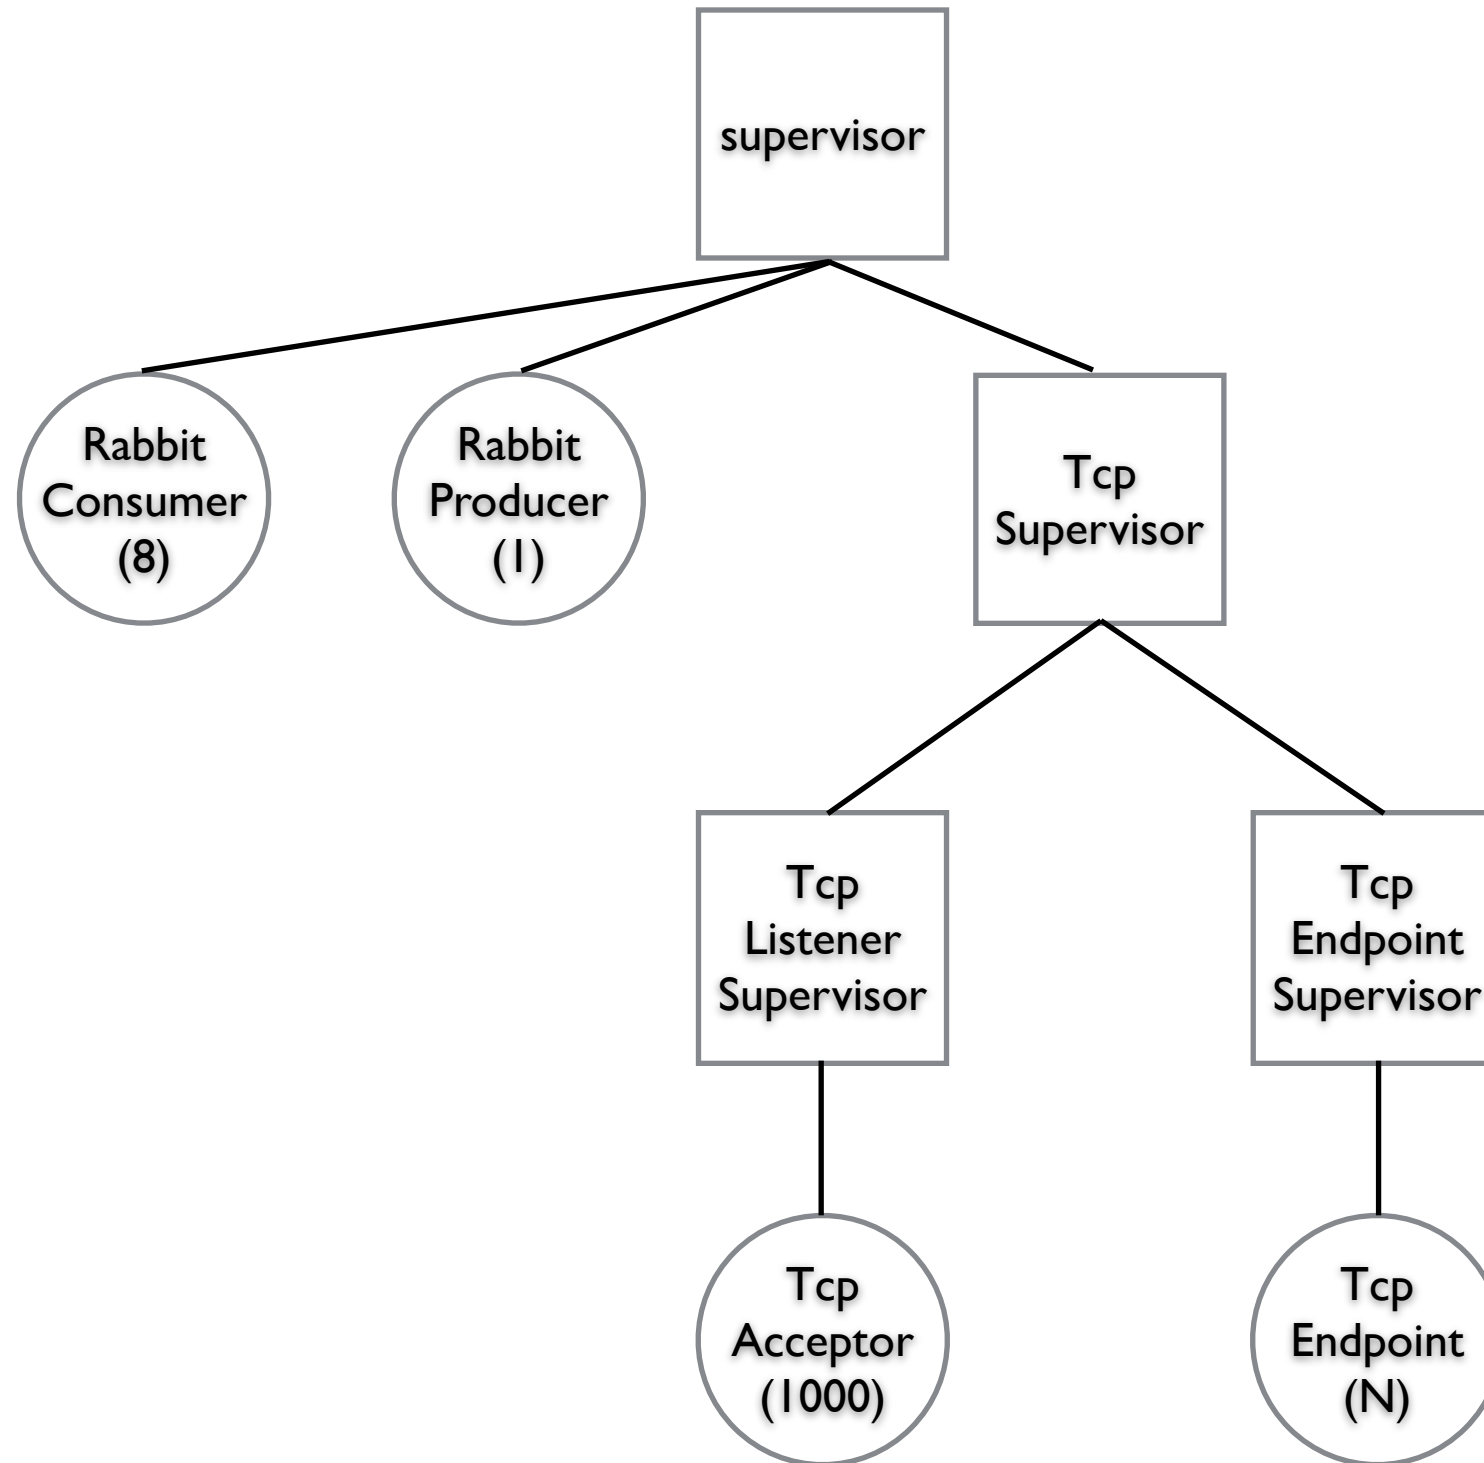
Processes

- Are not OS processes - much more lightweight (100000 on a single Erlang instance is not uncommon, millions per instance are doable)
- Can maintain state using recursive calls
- Communicate with other processes via messages and mailboxes
- Process messages in FIFO order

OTP - Open Telecom Platform

- Abstraction above simple processes that provides templates for actors
- Supervisors and Workers
 - If worker dies, supervisor is notified and can take action
- Defined, well understood pattern for handling state and message passing among processes

Supervisor Tree - Owsla



Code Walkthrough - Overview



An Accidental Acceptor Pool



Public Domain by <http://pdphoto.org/PictureDetail.php?mat=pdef&pg=5277>

Why do I need an Acceptor Pool?

- Single-threaded listen/accept loops artificially limit connections/second
- Acceptors can all share the listen socket, with only 1 being signaled for each connection
- Allows for concurrent acceptance of a very high number of connections/second (in our case, the goal was 1000/second)

Should I build my own like you did?

- NO
 - We didn't really know what an acceptor pool was when we started writing this code
 - Had we known, we would have found an existing library and used it instead.
- What libraries are out there?
 - Writing Erlang? Check out Ranch: <https://github.com/extend/ranch>
 - If you're writing Elixir, check out reagent: <https://github.com/meh/reagent>

Code Walkthrough - Tcp*



Taking Control of Restarts



Houston, We Have a Problem

- Load-testing tool would open 5-10k connections per instance to our server
- We'd run several instances of the load tool while testing
- Every time we stopped one of them, the entire TcpSupervisor tree would crash/restart, taking down all of the live connections

Restart Strategies

- `one_for_one`
 - If a child crashes, restart it
- `one_for_all`
 - if any child crashes, restart all of them
- `rest_for_one`
 - If a child crashes, it and all children started after it are restarted
- `simple_one_for_one`
 - like one-for-one, but designed for starting children dynamically.

Maximum Restart Frequency

- MaxT - the time window in which to measure restarts (in seconds) - default is 5
- MaxR - The maximum number of restarts that can occur within MaxT seconds before the supervisor will tear itself down - default is 5

Child Specifications: Restart

- permanent - always restart this child when it crashes (unless MaxR is exceeded)
- transient - restart if the child 'terminates abnormally'
- temporary - never restart, even if supervisor restart strategy tells you otherwise. These, therefore, aren't counted when looking at MaxR

Resources/Questions?

- <http://elixir-lang.org>
- <http://erlang.org>
- irc: #elixir-lang
- <http://jeetkundoug.wordpress.com/2014/01/13/elixir-and-the-internet-of-things>
- Me: @jeetkundoug
- My partners in crime on this project
 - @eymiha - Dave Anderson
 - @jvoegele - Jason Voegele